


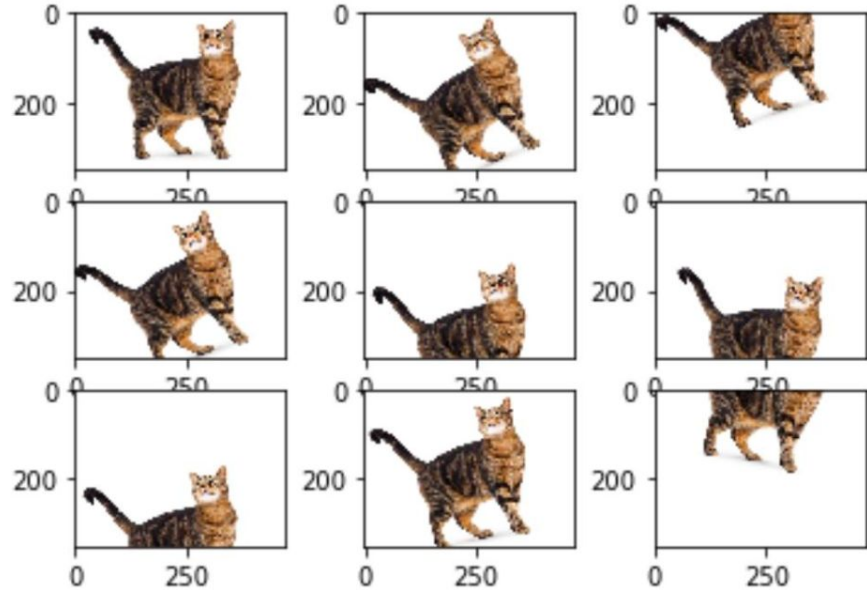
Convolutional Neural Networks II

Dr. Chelsea Parlett-Pelleriti

Convolutional Neural Networks



Data Augmentation



Data Augmentation

- Crop
- Flip
- Translation
- Rotation
- Zoom
- Contrast
- Brightness

Apply the augmentation stage to the batch of images.

```
plt.figure(figsize=(10, 10))
for images, _ in train_dataset.take(1):
    for i in range(9):
        augmented_images = data_augmentation(images)
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(augmented_images[0].numpy().astype("uint8"))
        plt.axis("off")
```

We can use `take(N)` to only sample `N` batches from the dataset. This is equivalent to inserting a `break` in the loop after the `N`th batch.

Display the first image in the output batch. For each of the nine iterations, this is a different augmentation of the same image.





Figure 8.10 Generating variations of a very good boy via random data augmentation

Residual Connections



Batch Normalization

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_1 \dots x_m\}$;

Parameters to be learned: γ, β

Output: $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

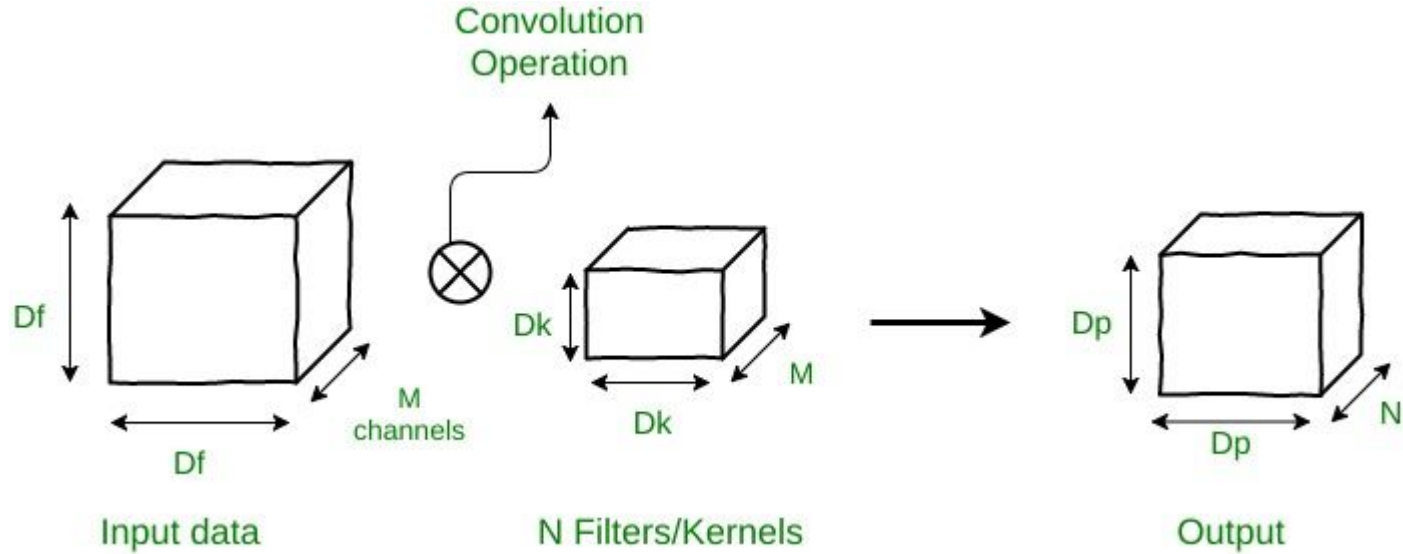
$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

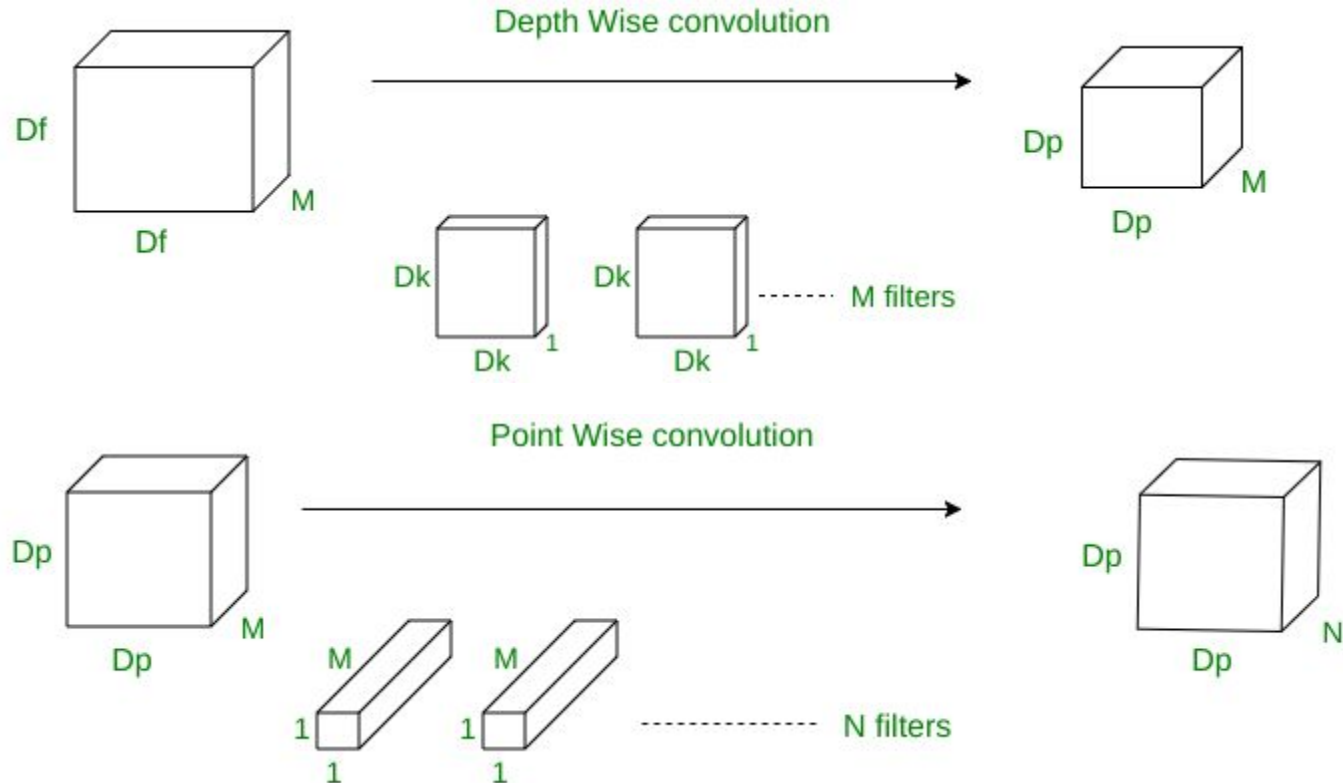
$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

Algorithm 1: Batch Normalizing Transform, applied to activation x over a mini-batch.

Depthwise Separable Convolutions




Depthwise Separable Convolutions



Depthwise Separable Convolutions

Type of Convolution	Complexity
Standard	$N \times D_p^2 \times D_g^2 \times M$
Depth wise separable	$M \times D_p^2 \times (D_k^2 + N)$

Visualizing Convnets




Interpretable AI

Visualizing ConvNets


- Visualizing Layer Activations
- Visualizing Filters
- Visualizing Class Activation Heatmaps

Visualizing Layer Activations

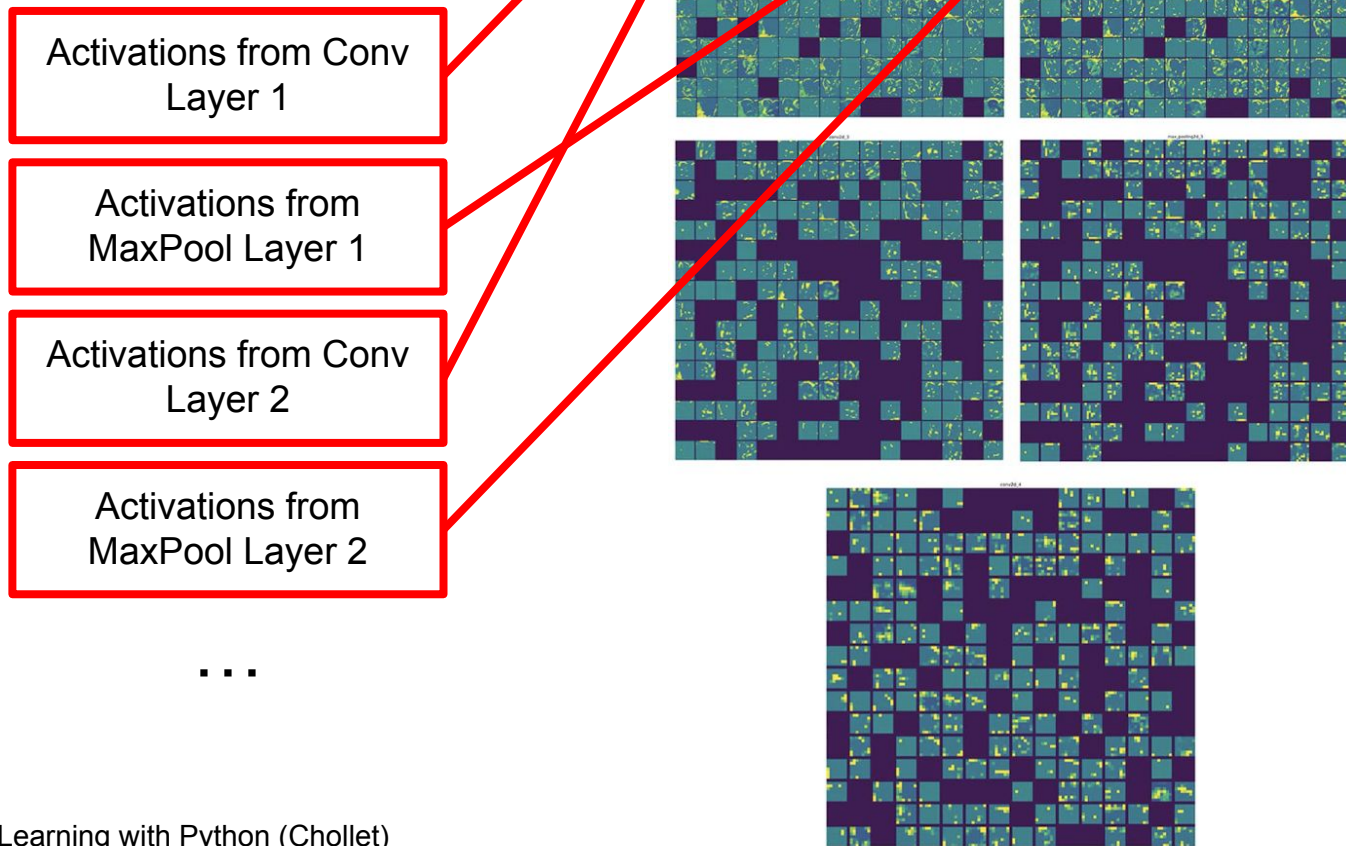


Visualizing Layer Activations

1. Take a test image
2. Feed it through each layer and save the output/activations
3. Plot



Visualizing Layer Activations



Visualizing Layer Activations

You can see:

- The sparsity of higher layers (due to ReLu)
- The increasing abstraction of higher layers

Visualizing Filters

- What image would the filter *maximally* respond to? What activates it the most?
- Start with a **blank image** and use **gradient descent** to change it until the filter responds maximally

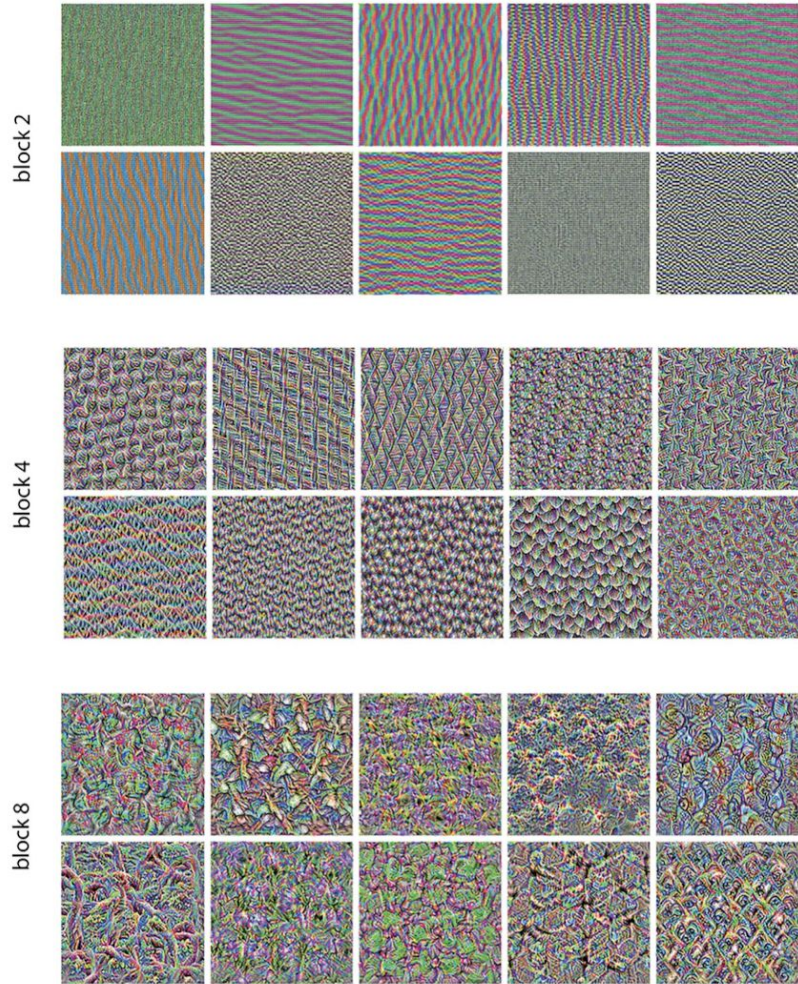



Figure 9.17 Some filter patterns for layers block2_sepconv1, block4_sepconv1, and block8_sepconv1


Visualizing Filters

- Filter seems to respond to horizontal lines



Visualizing Class Activation Heatmaps

Grad-CAM:
creates
heatmaps of how
intensely the
input images
activates the
class



Visualizing Class Activation Heatmaps

How *intensely* the input image **activated different channels** in the last layer


- Send image through network and get activations of final conv layer (before classification layer)

How *important* each **channel** is with regard to that class

- calculate the gradient for a specific class with respect to the activations of the final conv layer


Visualizing Class Activation Heatmaps

Heatmap of how intensely the input images activates the class



Visualizing Class Activation Heatmaps

Heatmap of how
intense
input



NOTE: We usually create these heatmaps for the top/predicted class but we CAN create it for any class




Visualizing Class Activation Heatmaps



Predicted: **Wolf**
True: **Wolf**



Predicted: **Husky**
True: **Husky**



Predicted: **Husky**
True: **Husky**



Predicted: **Wolf**
True: **Wolf**



Predicted: **Wolf**
True: **Wolf**




Predicted: **Wolf**
True: **Wolf**




Predicted: **Husky**
True: **Wolf**



Predicted: **Wolf**
True: **Wolf**



Predicted: **Wolf**
True: **Husky**




Predicted: **Husky**
True: **Husky**



CAM would be helpful in situations like this!

Bonus: Occlusion

Mask part of the image before feeding to CNN,
check how much predicted probabilities change



$$P(\text{elephant}) = 0.95$$



$$P(\text{elephant}) = 0.75$$

Bonus: Tricking Convnets



x

“panda”

57.7% confidence

+ .007 ×



$\text{sign}(\nabla_x J(\theta, x, y))$

“nematode”

8.2% confidence

=



$x + \epsilon \text{sign}(\nabla_x J(\theta, x, y))$
“gibbon”
99.3 % confidence

Types of Computer Vision Tasks

- Image Classification (done)
- Image Segmentation
- Object Detection
- Inpainting

Image Segmentation



Figure 9.2 Semantic segmentation vs. instance segmentation

Image Segmentation


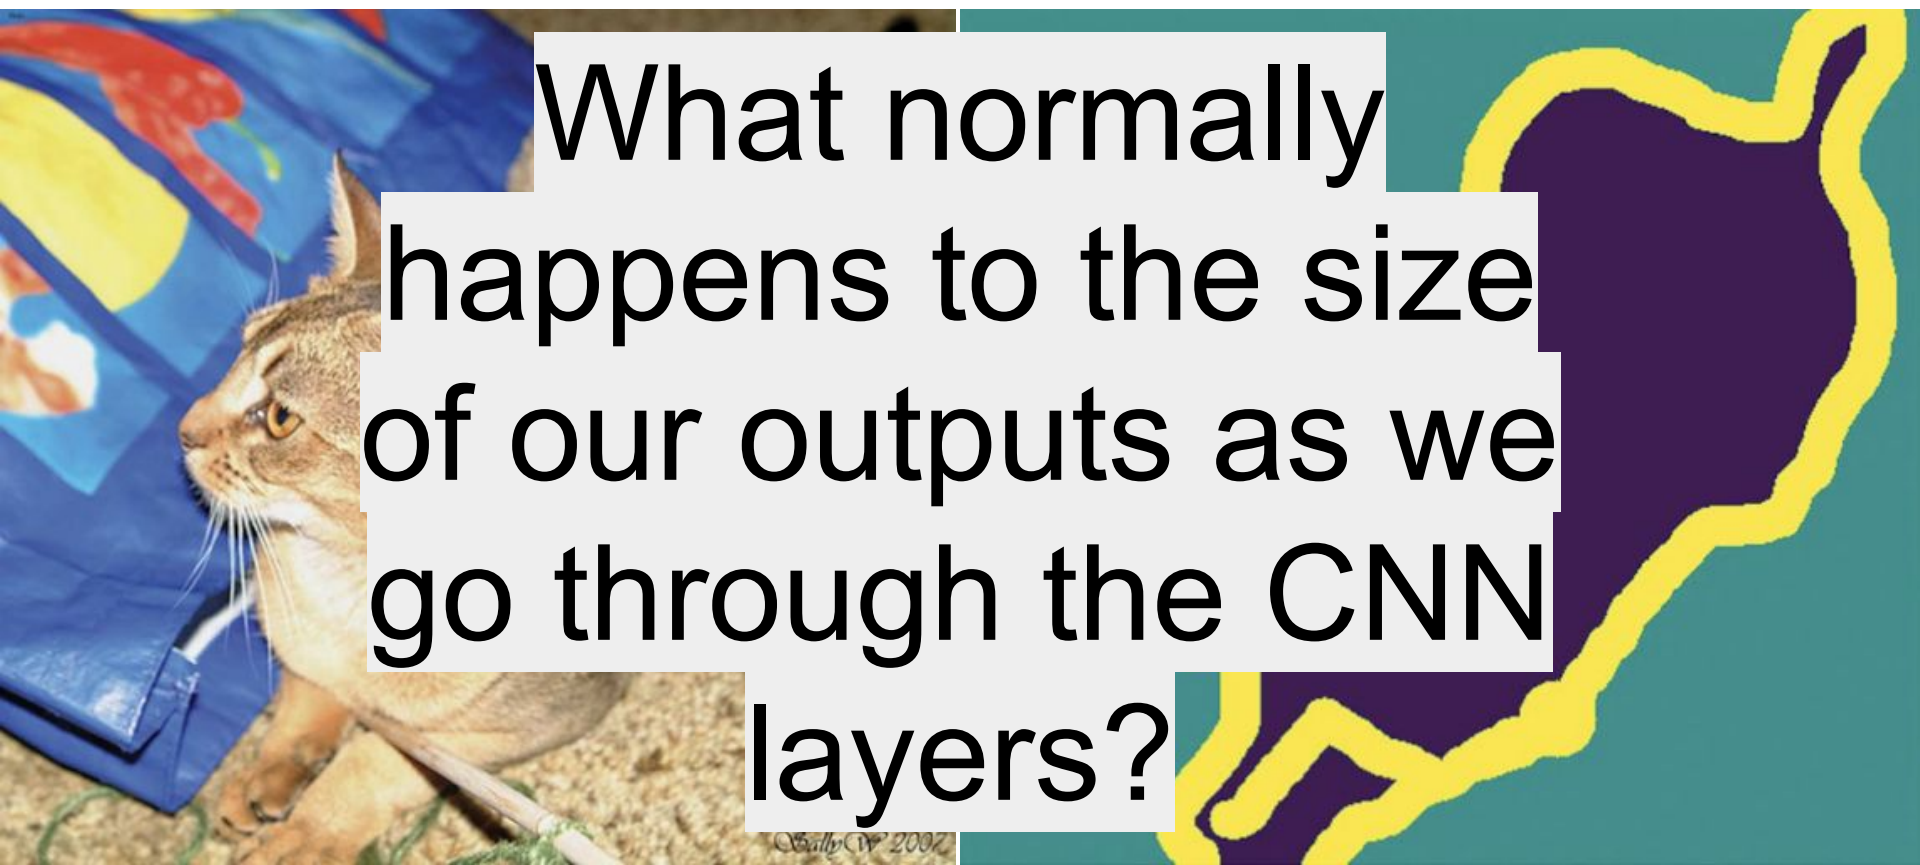



Image Segmentation



What normally happens to the size of our outputs as we go through the CNN layers?

Object Detection

Object detection



Neural Inpainting




Image from: Altinel, Fazil & Ozay, Mete & Okatani, Takayuki. (2018). Deep Structured Energy-Based Image Inpainting.

Pre-trained Models

