

Generative Models II

Dr. Parlett-Pelleriti

Outline

- Generative Adversarial Networks Overview
- Training GANs
- Issues with Training GANs
- Conditional GANs

Generative Adversarial Networks

Generative **Adversarial** Networks



Generator



From:
<https://www.youtube.com/watch?v=JBIm4wnjNMY>

Discriminator



From:
<https://www.youtube.com/watch?v=JBIm4wnjNMY>

Adversarial Training



From:
<https://www.youtube.com/watch?v=JBIm4wnjNMY>

Adversarial Training



x

“panda”

57.7% confidence

$+ .007 \times$



$\text{sign}(\nabla_x J(\theta, x, y))$

“nematode”

8.2% confidence

$=$



$x + \epsilon \text{sign}(\nabla_x J(\theta, x, y))$

“gibbon”

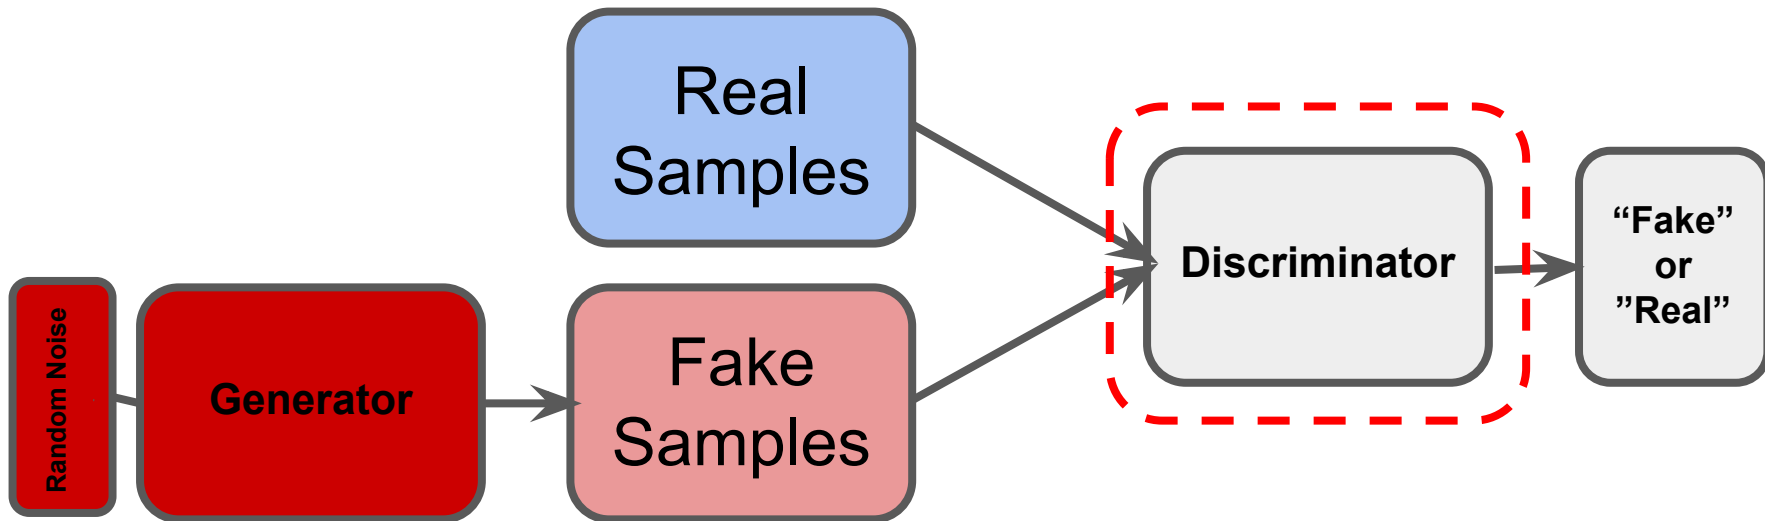
99.3 % confidence

Adversarial Training

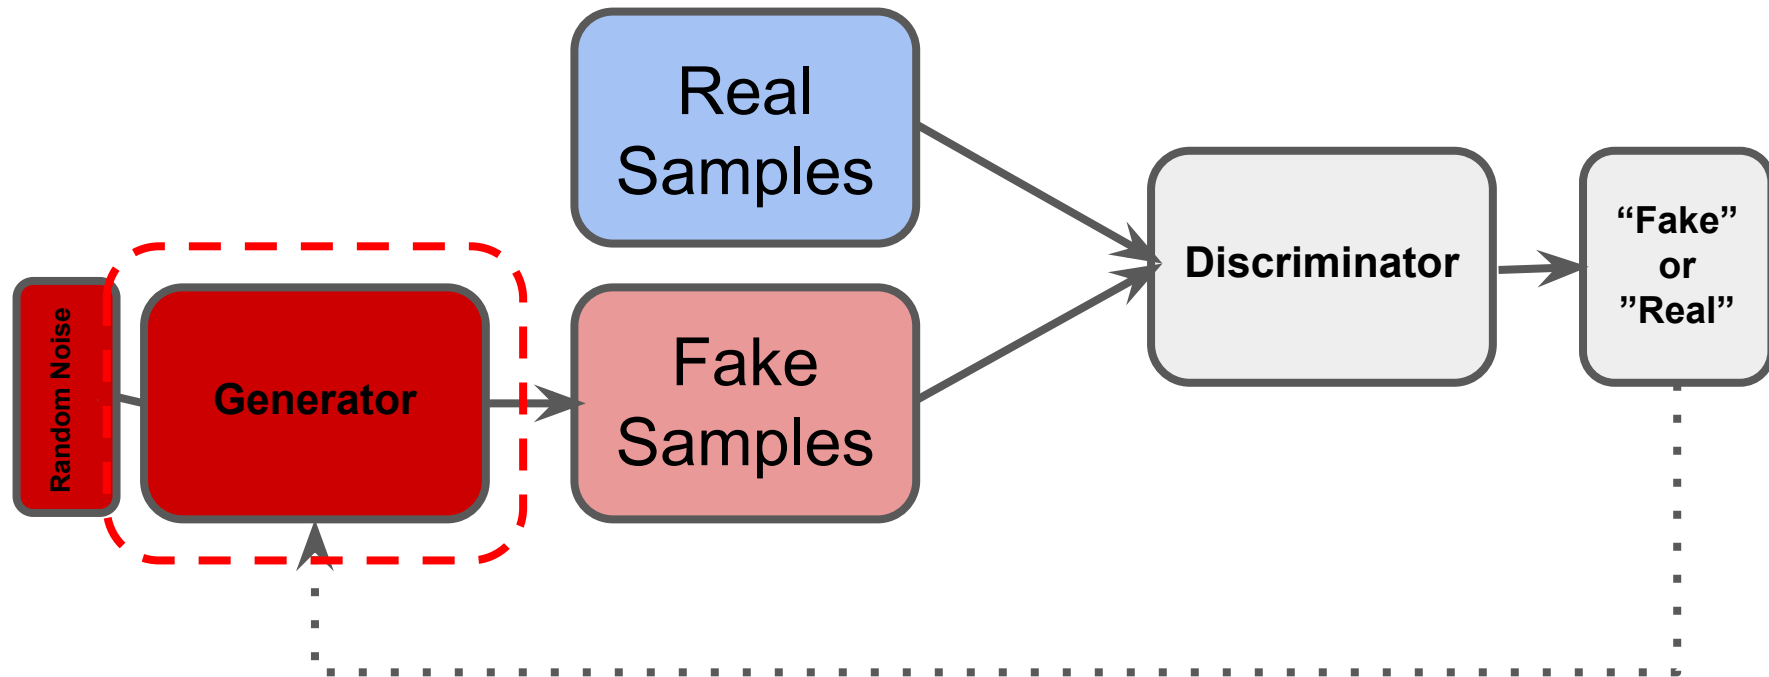


ZeroGPT

The Discriminator




The Generator



Training GANs

1. Train the **discriminator** holding **generator** constant
2. Train the **generator** holding the **discriminator** constant
3. **Repeat** until convergence

GAN Loss Functions



GAN Loss Functions

$$E_x [\log(D(x))] + E_z [\log(1 - D(G(z)))]$$

Discriminator's
estimate that real
sample is real

GAN Loss Functions

$$E_x [\log(D(x))] + E_z [\log(1 - D(G(z)))]$$

Discriminator's
estimate that fake
sample is real

GAN Loss Functions

Discriminator wants to maximize

$$E_x [\log(D(x))] + E_z [\log(1 - D(G(z)))]$$

Discriminator's
estimate that real
sample is real

Discriminator's
estimate that fake
sample is real

GAN Loss Functions

Generator wants to minimize

$$E_x [\log(D(x))] + E_z [\log(1 - D(G(z)))]$$

Discriminator's
estimate that real
sample is real

Discriminator's
estimate that fake
sample is real

Training Loop

```
for i in range(training_iteration):  
    # hold generator constant  
    for k in range(discrim_steps):  
        x = sample_real_inputs(num = m)  
        gz1 = generate_fake_inputs(num = m)  
        discrim_params = gradient_ascent(x, gz)  
    # hold discriminator constant  
    gz2 = generate_fake_inputs(num = m)  
    gener_params = gradient_descent(gz2)
```

Training Loop

```
for i in range(training_iteration):  
    # hold generator constant  
    for k in range(discrim_steps):  
        x = sample_real_inputs(num = m)  
        gz1 = generate_fake_inputs(num = m)  
        discrim_params = gradient_ascent(x, gz)  
    # hold discriminator constant  
    gz2 = generate_fake_inputs(num = m)  
    gener_params = gradient_descent(gz2)
```


Training Loop

```
for i in range(training_iteration):  
    # hold generator constant  
    for k in range(discrim_steps):  
        x = sample_real_inputs(num = m)  
        gz1 = generate_fake_inputs(num = m)  
        discrim_params = gradient_ascent(x, gz)  
        # hold discriminator constant  
        gz2 = generate_fake_inputs(num = m)  
        gener_params = gradient_descent(gz2)
```

Training GANs


1. Train the **discriminator** holding **generator** fixed
2. Train the **generator** holding the **discriminator** fixed
3. **Repeat** until convergence

What if the
Discriminator is too
good?



*In the minimax game, the discriminator minimizes a cross-entropy, but the generator maximizes the same cross-entropy. This is unfortunate for the generator, because when the discriminator successfully rejects generator samples with high confidence, the **generator's gradient vanishes**.*

-Ian Goodfellow (<https://arxiv.org/pdf/1701.00160.pdf>)



Non-Saturating GAN Loss

Generator wants to maximize

$$-\log(D(G(z)))$$

instead of minimizing


$$\log(1 - D(G(z)))$$

Mode Collapse




Image from: <https://neptune.ai/blog/gan-loss-functions>

Mode Collapse



Convergence



Sometimes the two players eventually reach an equilibrium, but in other scenarios they repeatedly undo each others' progress without arriving anywhere useful.

-Ian Goodfellow

(<https://arxiv.org/pdf/1701.00160.pdf>)

Wasserstein GANs



Discriminator Maximizes:

$$D(x) - D(G(z))$$

Generator Maximizes:

$$D(G(z))$$

Wasserstein GANs



“The Critic” Maximizes:

$$D(x) - D(G(z))$$

Generator Maximizes:

$$D(G(z))$$

Conditional GANs

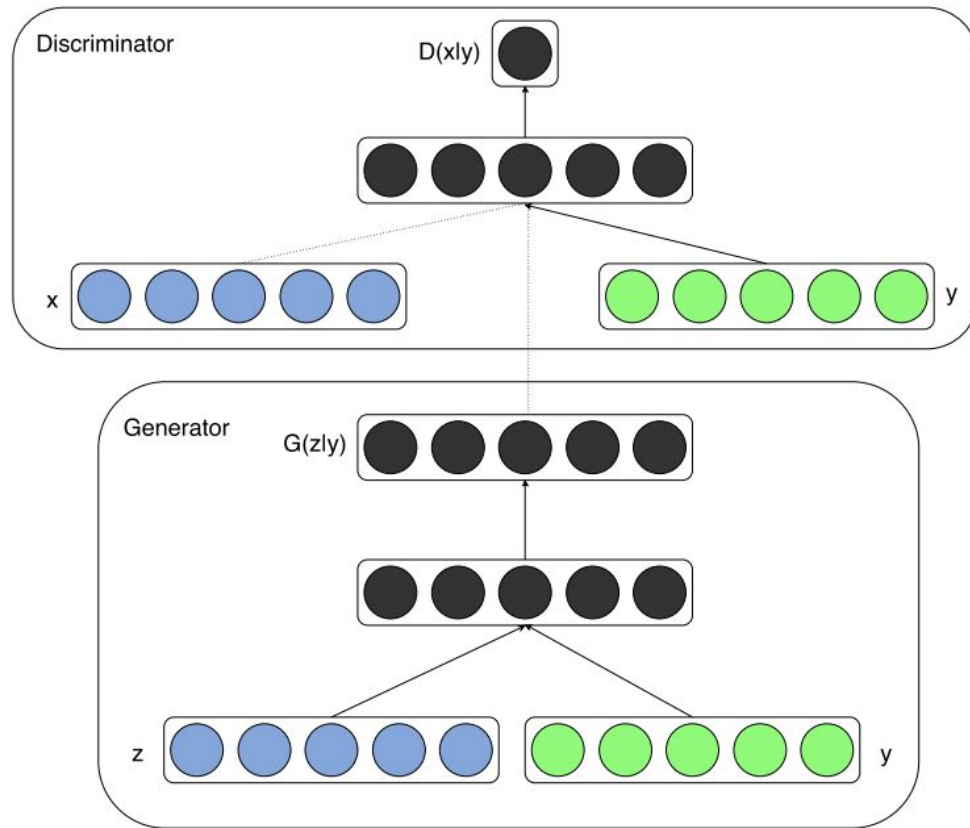


Figure 1: Conditional adversarial net

Conditional GAN

