Transformers I

Dr. Parlett-Pelleriti

Transformers

Outline

- Encoder/Decoder Structure
- Transformer Intro
- Word Embeddings
- Positional Encoding

Sequence Models

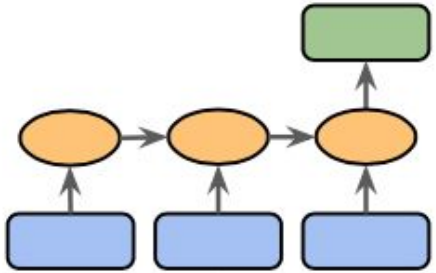

many-to-one

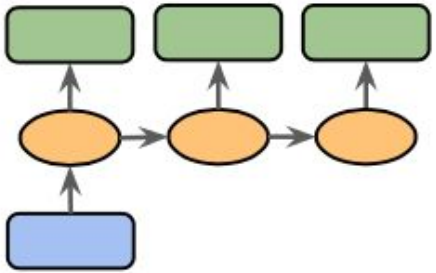

one-to-many

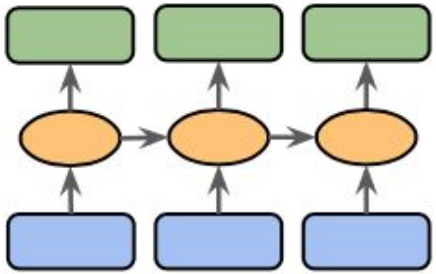

many-to-many

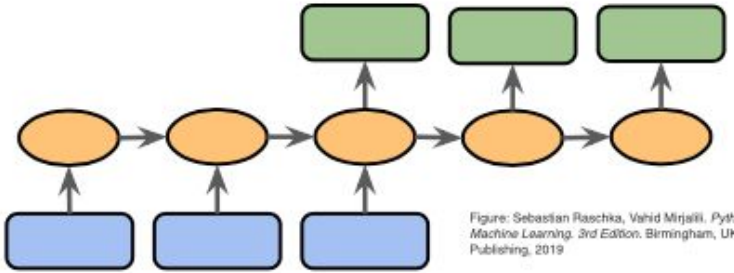

Figure: Sebastian Raschka, Vahid Mirjalili. *Python Machine Learning. 3rd Edition*. Birmingham, UK: Packt Publishing, 2019

many-to-many

Seq-to-Seq Models

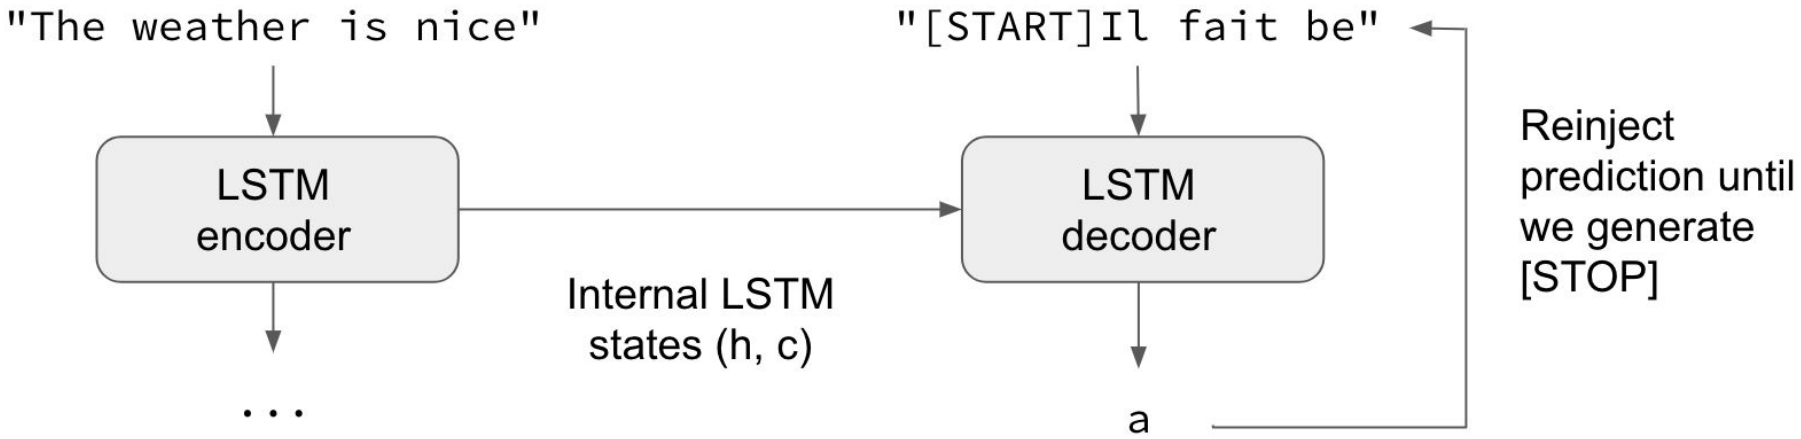

Image from: https://blog.keras.io/a-ten-minute-introduction-to-sequence-to-sequence-learning-in-keras.html

Autoencoder Architecture

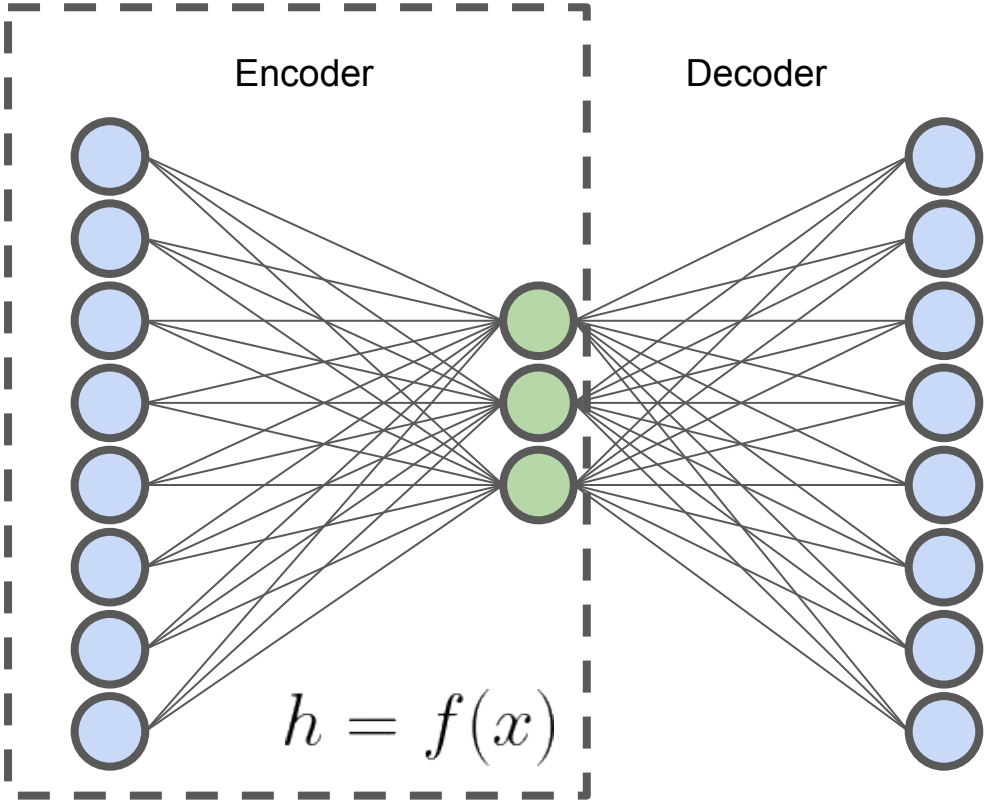

Autoencoder Architecture

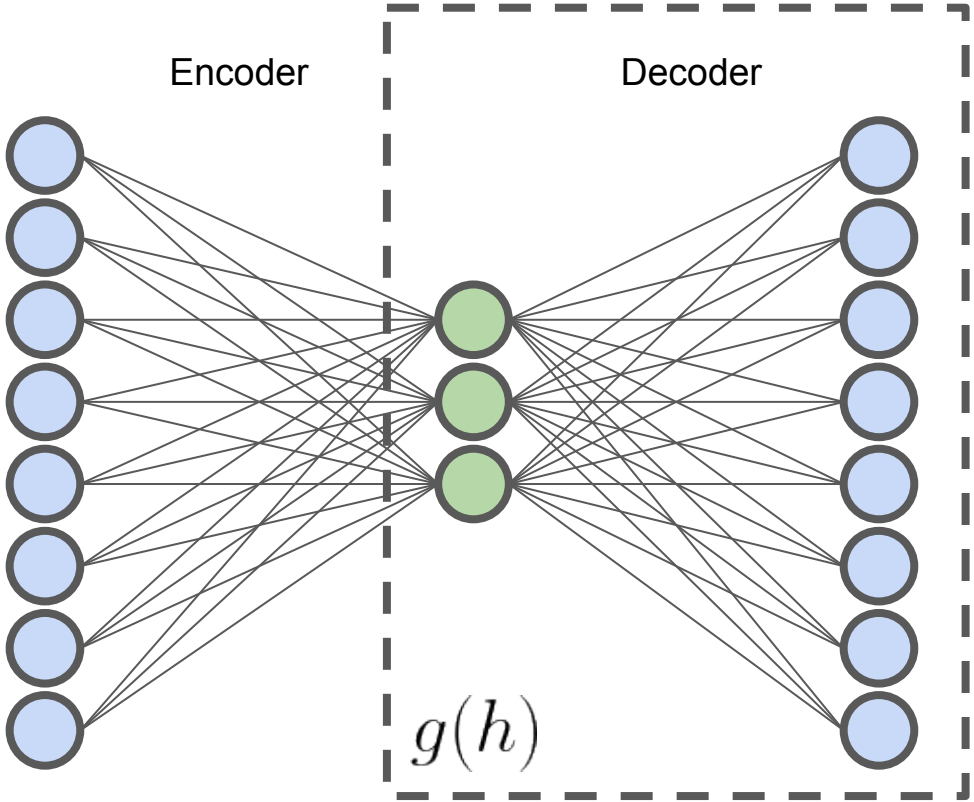

Autoencoder Architecture

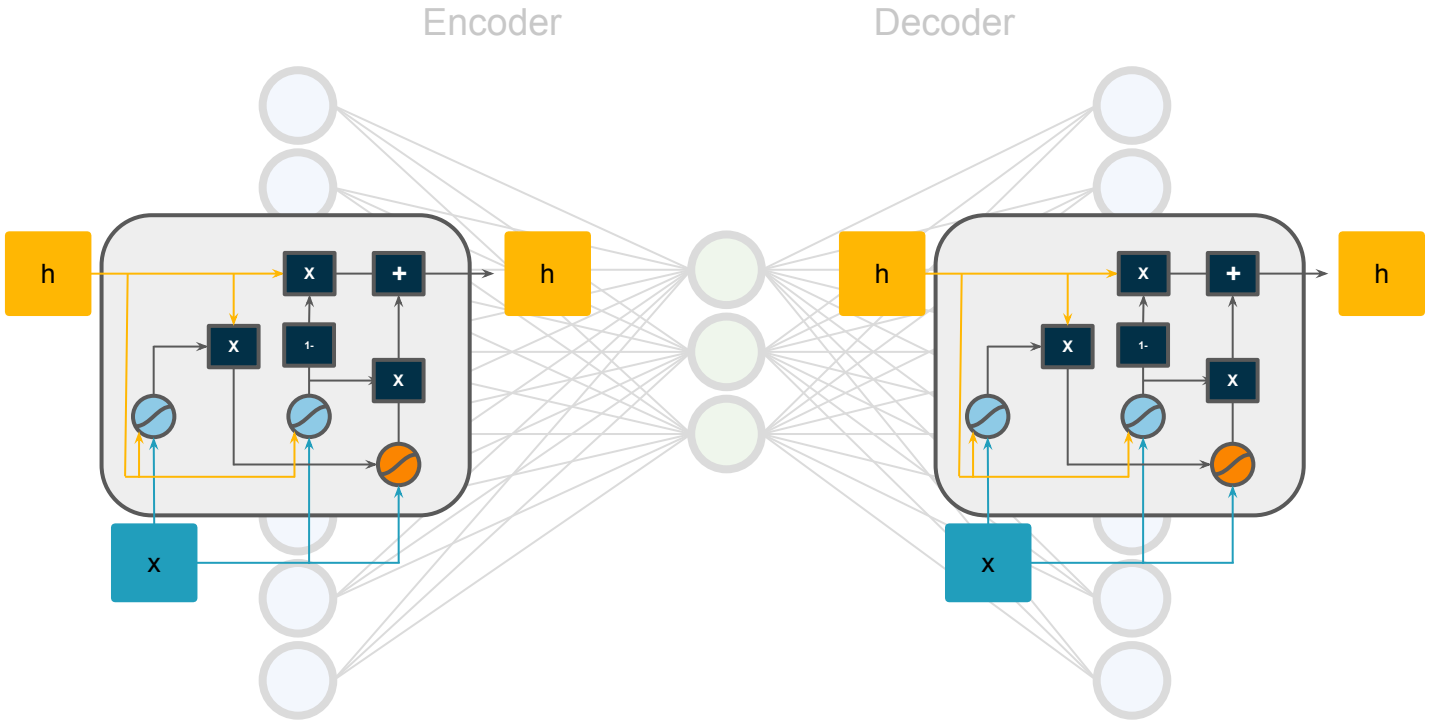

Cho et. al
2014

GRU Machine Translation

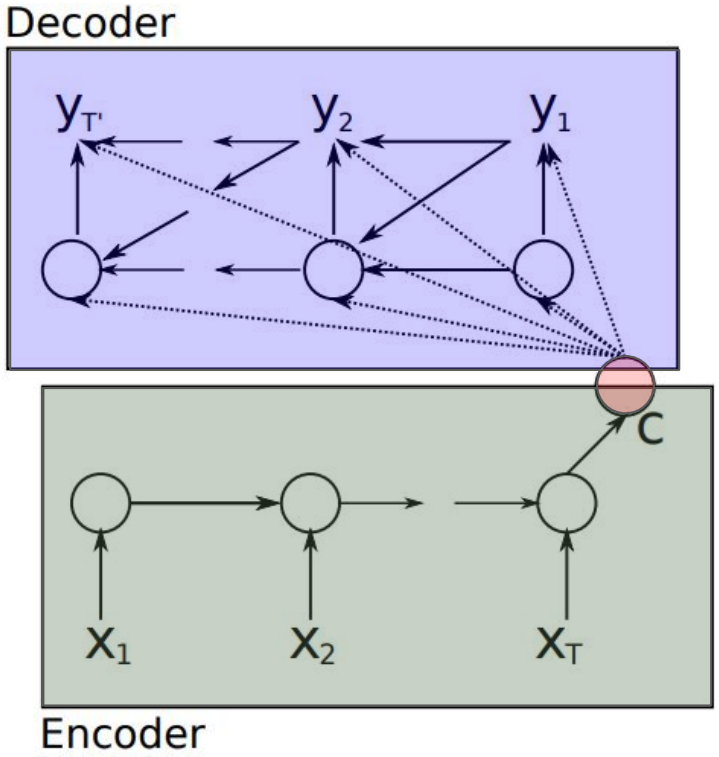

Cho et. al
2014

General Encoder/Decoder Model

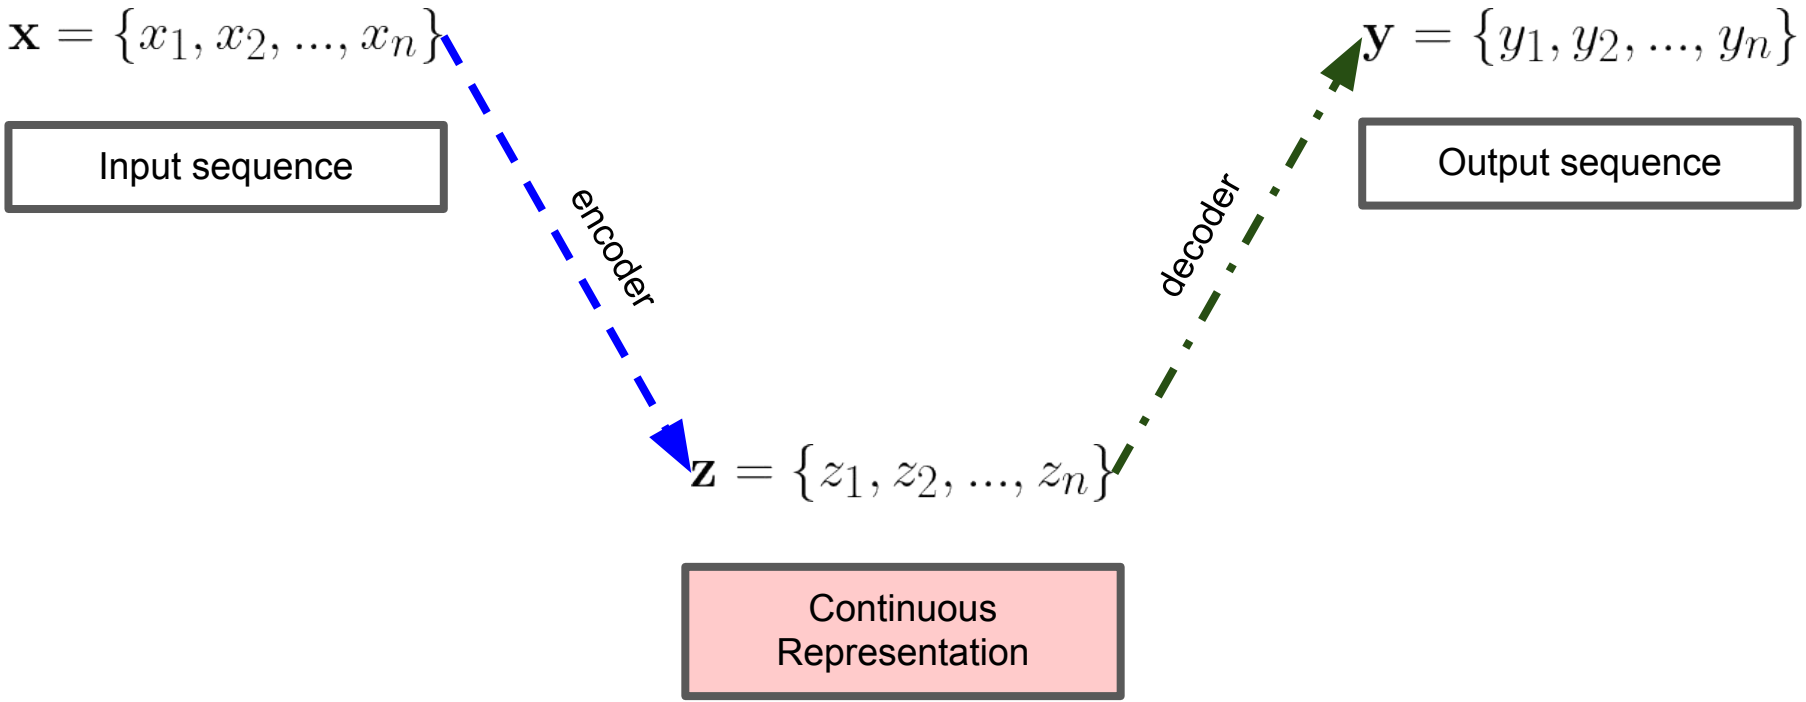

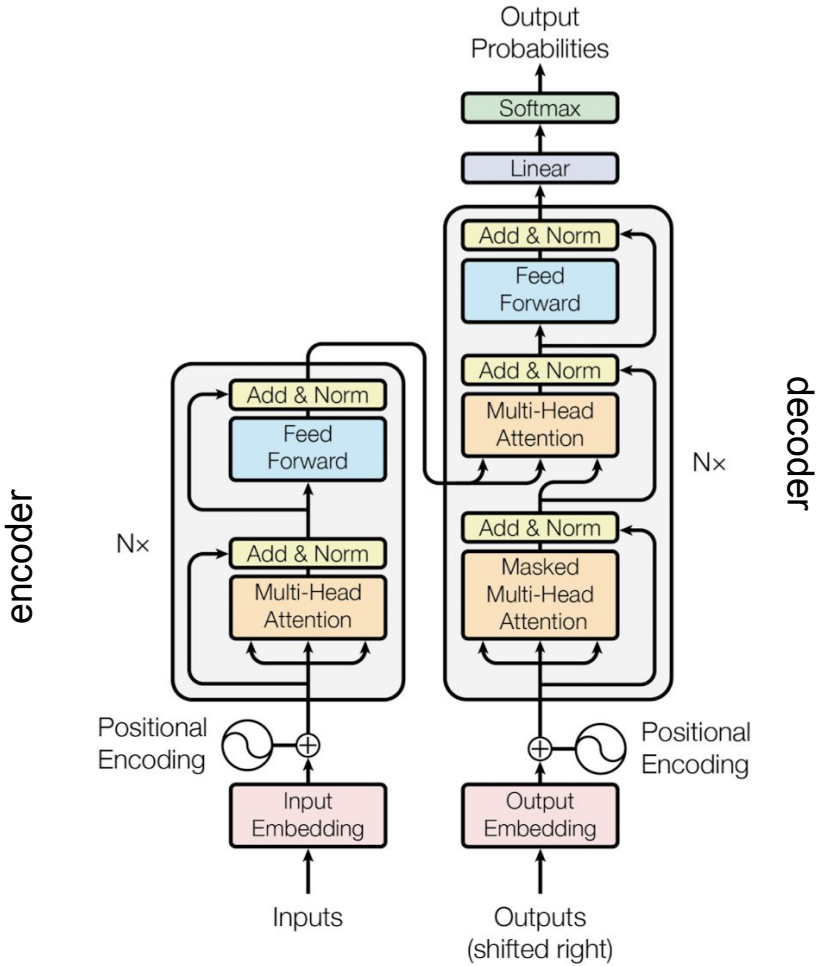

Figure 1: The Transformer - model architecture.

Image from: Attention is all you Need (Vaswani et. al 2017)

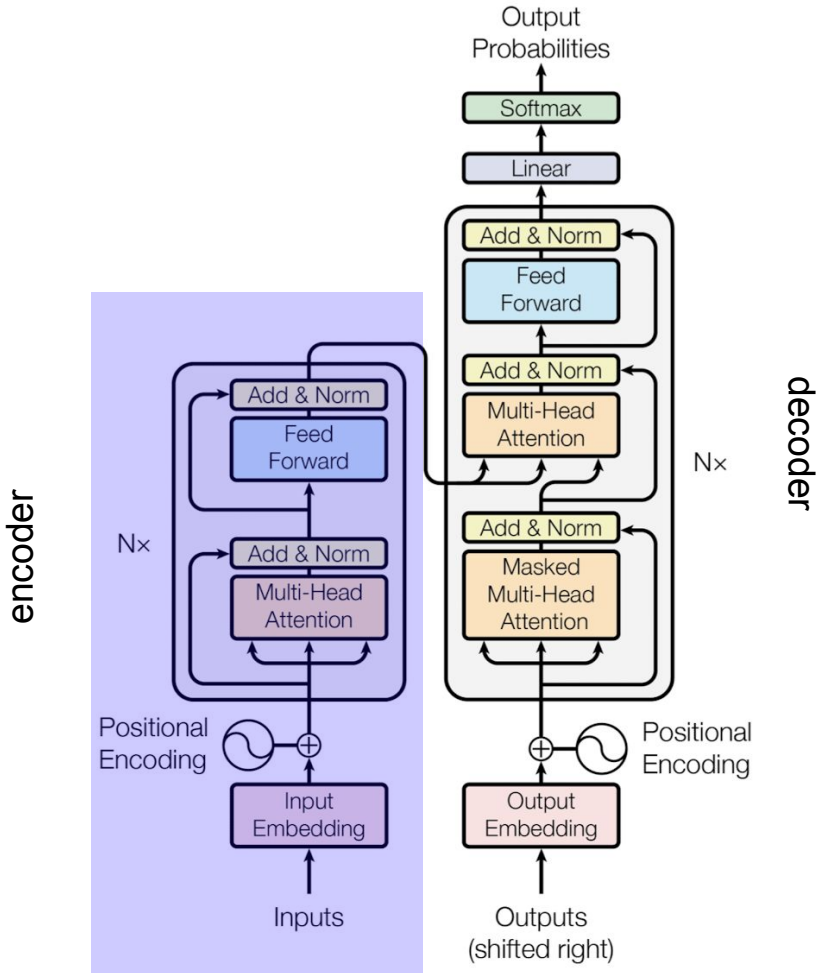

Figure 1: The Transformer - model architecture.

Image from: Attention is all you Need (Vaswani et. al 2017)

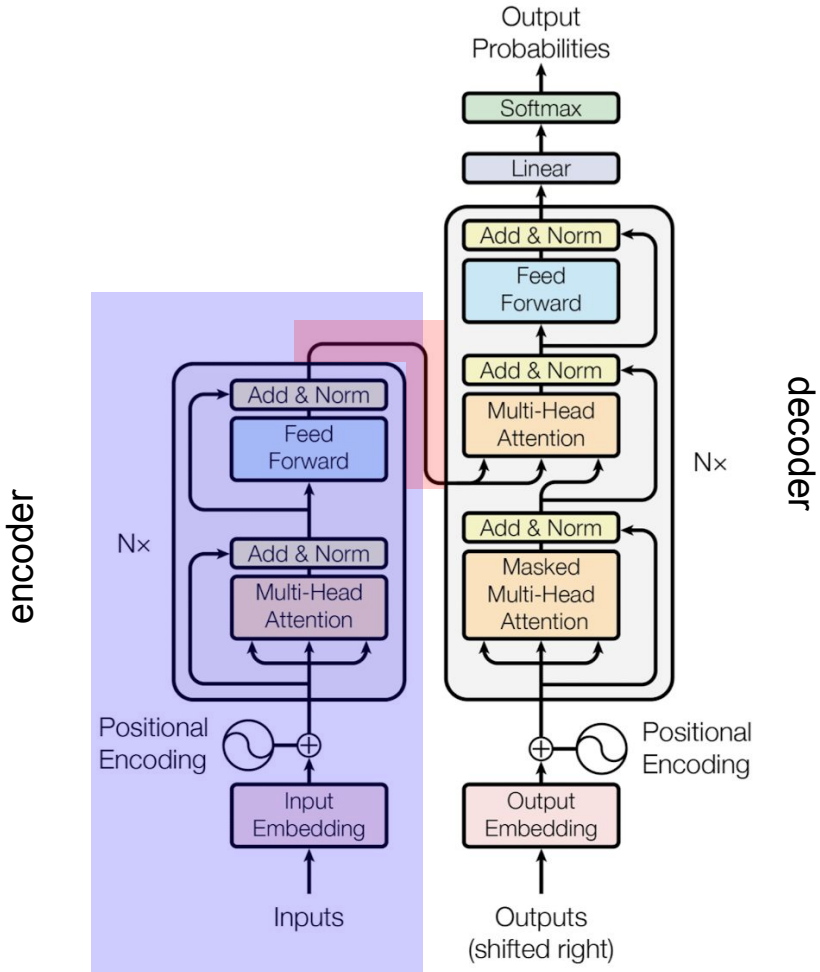

Figure 1: The Transformer - model architecture.

Image from: Attention is all you Need (Vaswani et. al 2017)

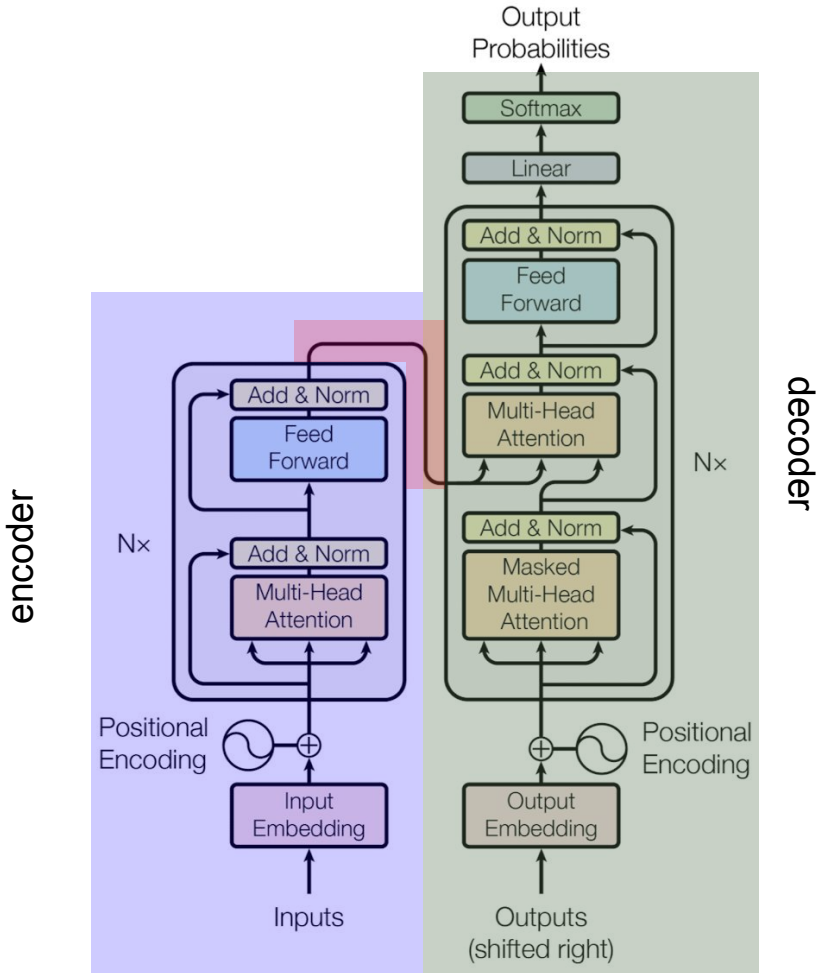

Figure 1: The Transformer - model architecture.

Image from: Attention is all you Need (Vaswani et. al 2017)

Attention Is All You Need

Ashish Vaswani* Google Brain avaswani@google.com

Noam Shazeer* Google Brain noam@google.com

Niki Parmar* Google Research nikip@google.com

Jakob Uszkoreit* Google Research usz@google.com

Llion Jones* Google Research llion@google.com

Aidan N. Gomez* † University of Toronto aidan@cs.toronto.edu

Łukasz Kaiser* Google Brain lukaszkaiser@google.com

Illia Polosukhin* ‡ illia.polosukhin@gmail.com

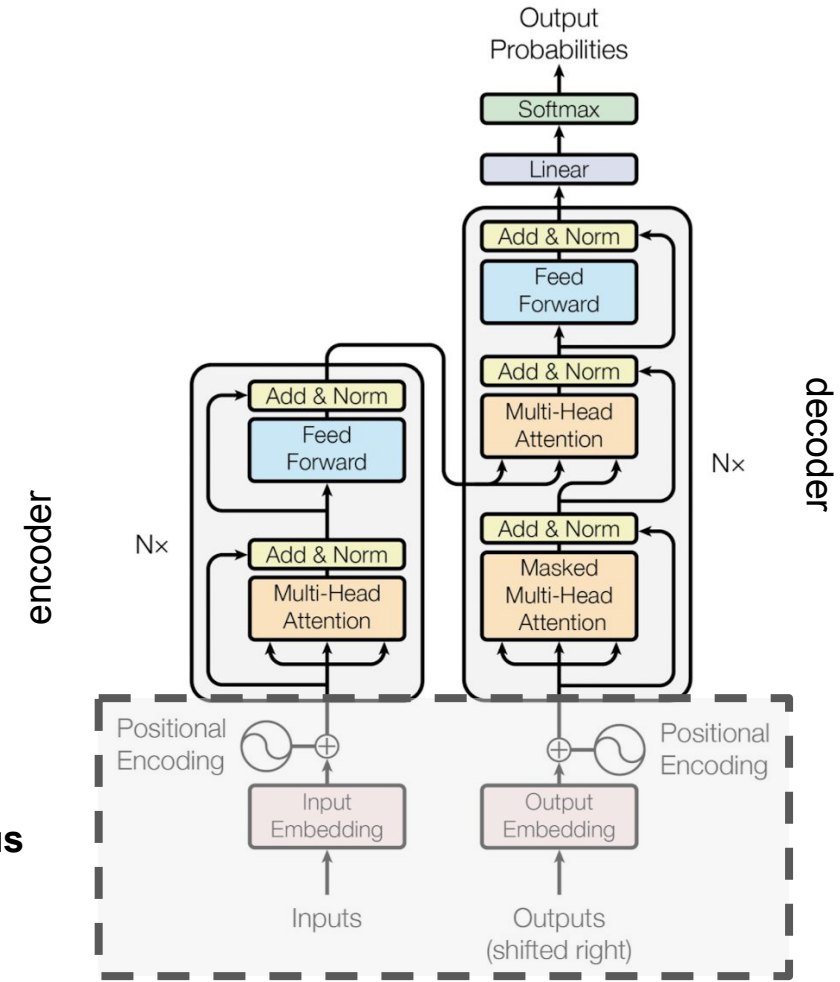

Figure 1: The Transformer - model architecture.

Today's Focus

Image from: Attention is all you Need (Vaswani et. al 2017)

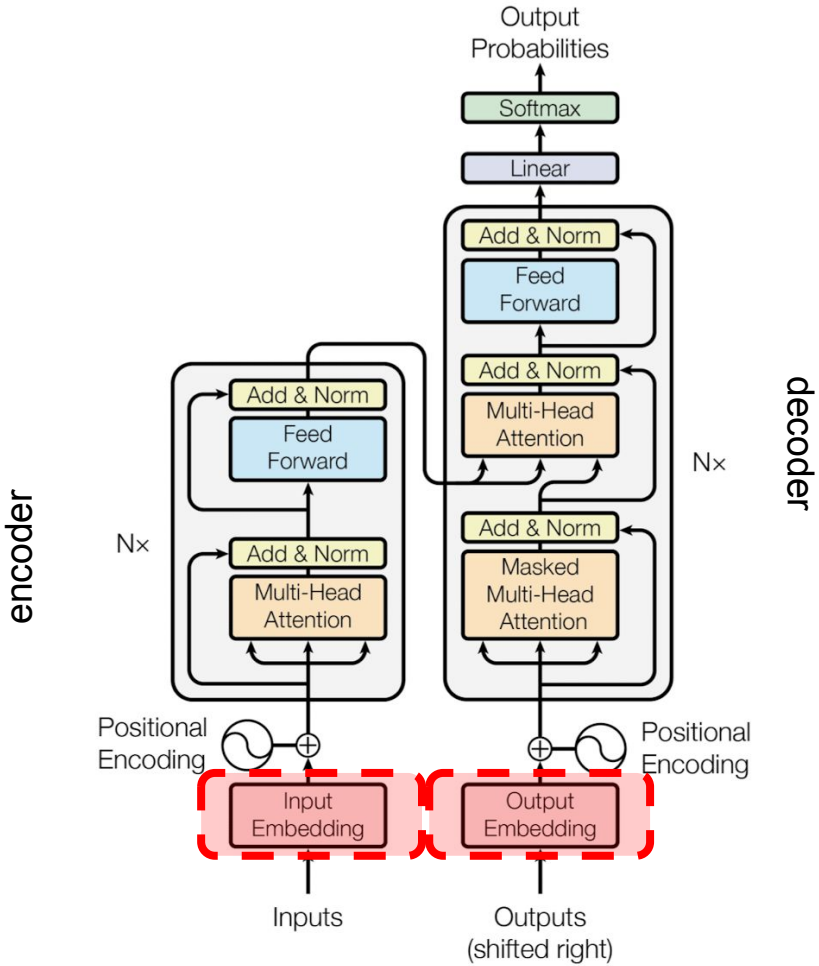

Figure 1: The Transformer - model architecture.

Image from: Attention is all you Need (Vaswani et. al 2017)

Word Embeddings

$$\begin{bmatrix} 0.1 \\ 0.2 \\ 0.4 \\ 0.2 \\ 0.6 \\ 0.1 \end{bmatrix}$$

gorgeous

$$\begin{bmatrix} 0.2 \\ 0.4 \\ 0.5 \\ 0.9 \\ 0.3 \\ 0.1 \end{bmatrix}$$

python

Word Embeddings

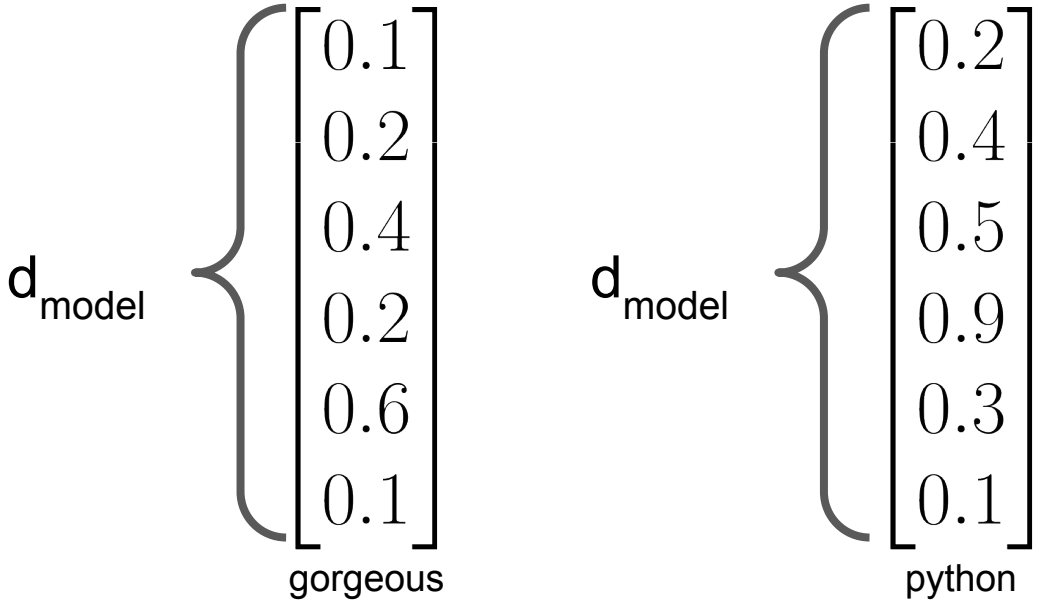

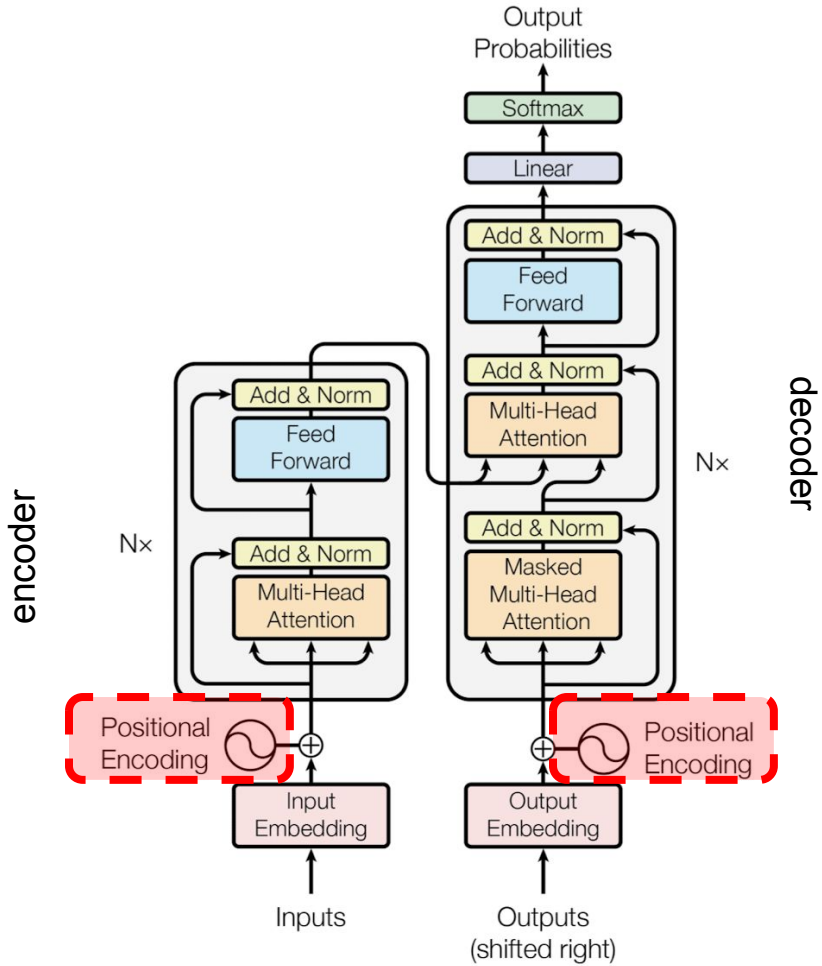

Figure 1: The Transformer - model architecture.

Image from: Attention is all you Need (Vaswani et. al 2017)

Positional Encoding

the | cat | ate | the | bat

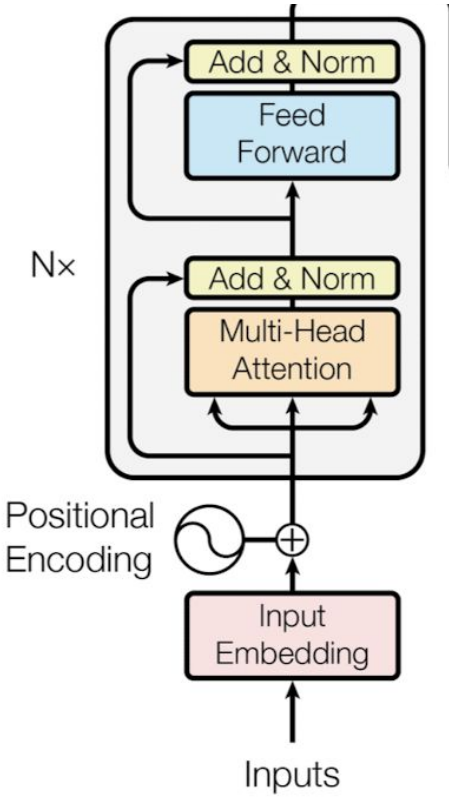

Positional Encoding

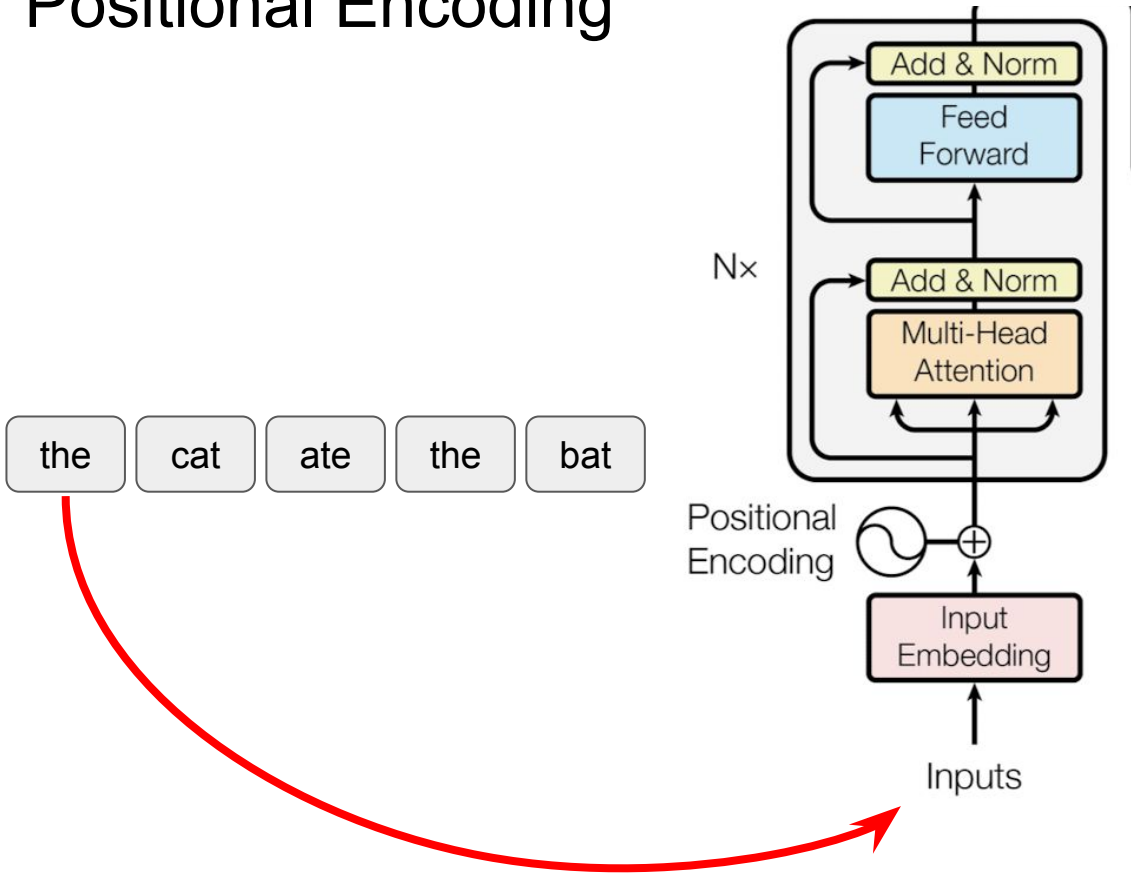

Positional Encoding

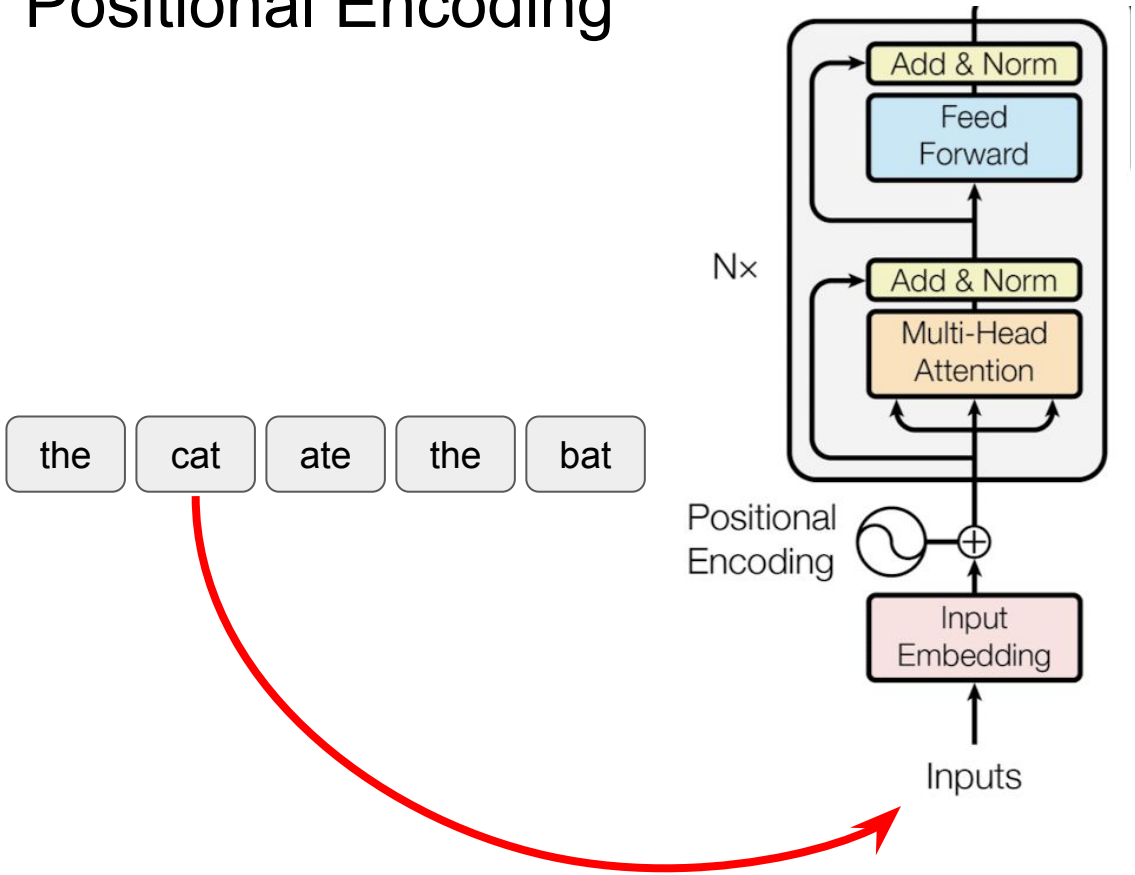

Positional Encoding

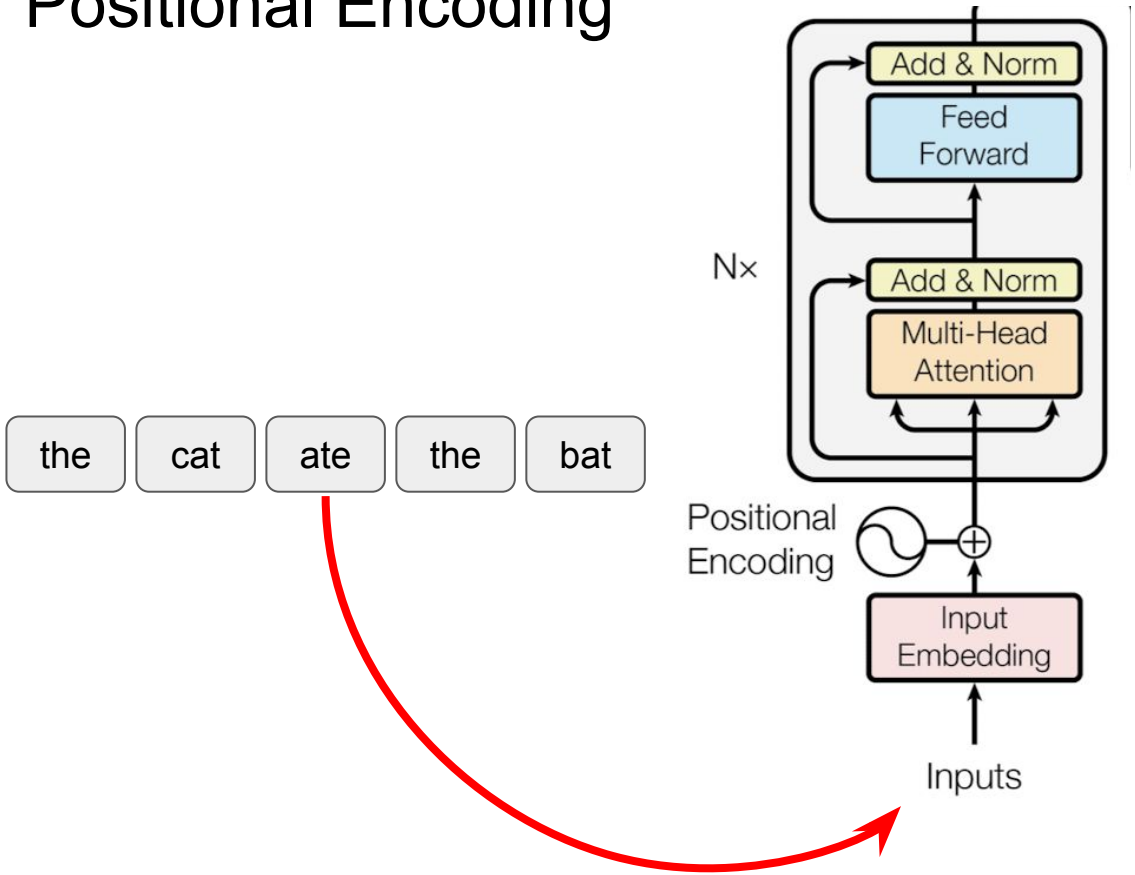

Positional Encoding

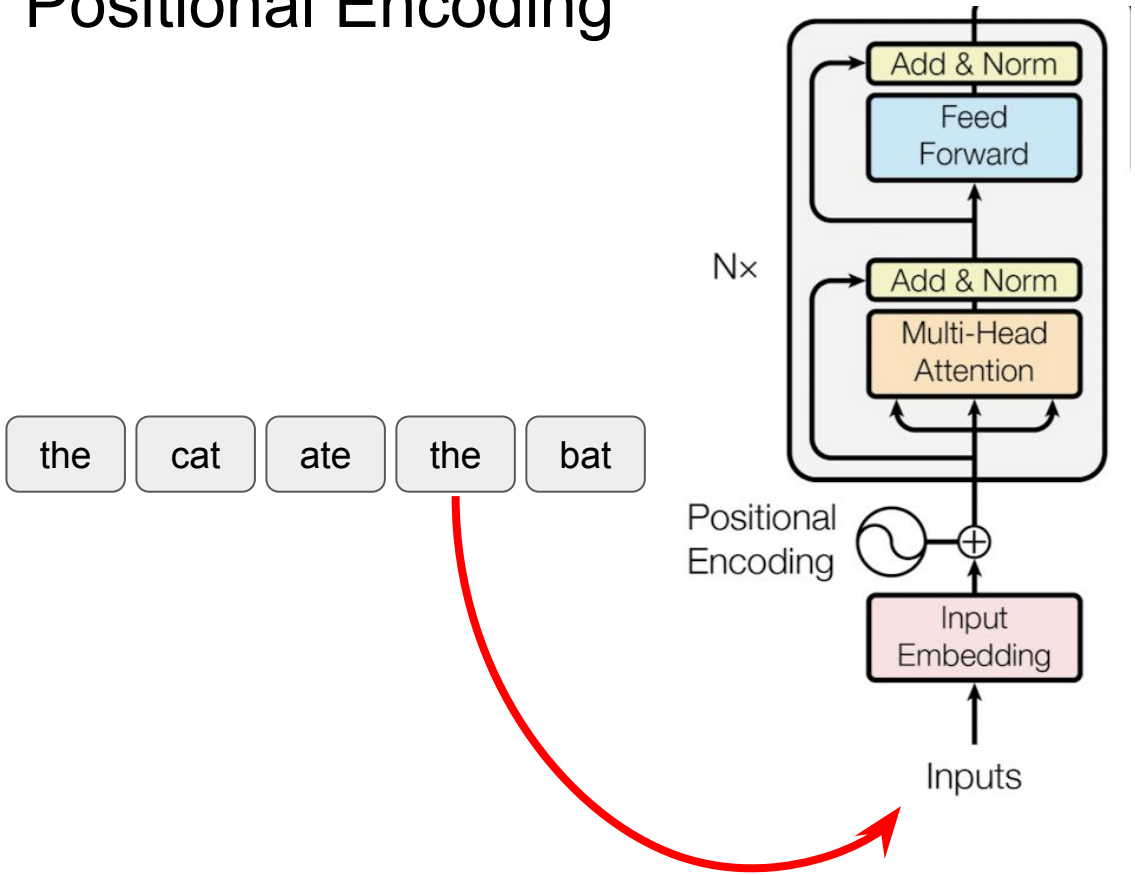

Positional Encoding

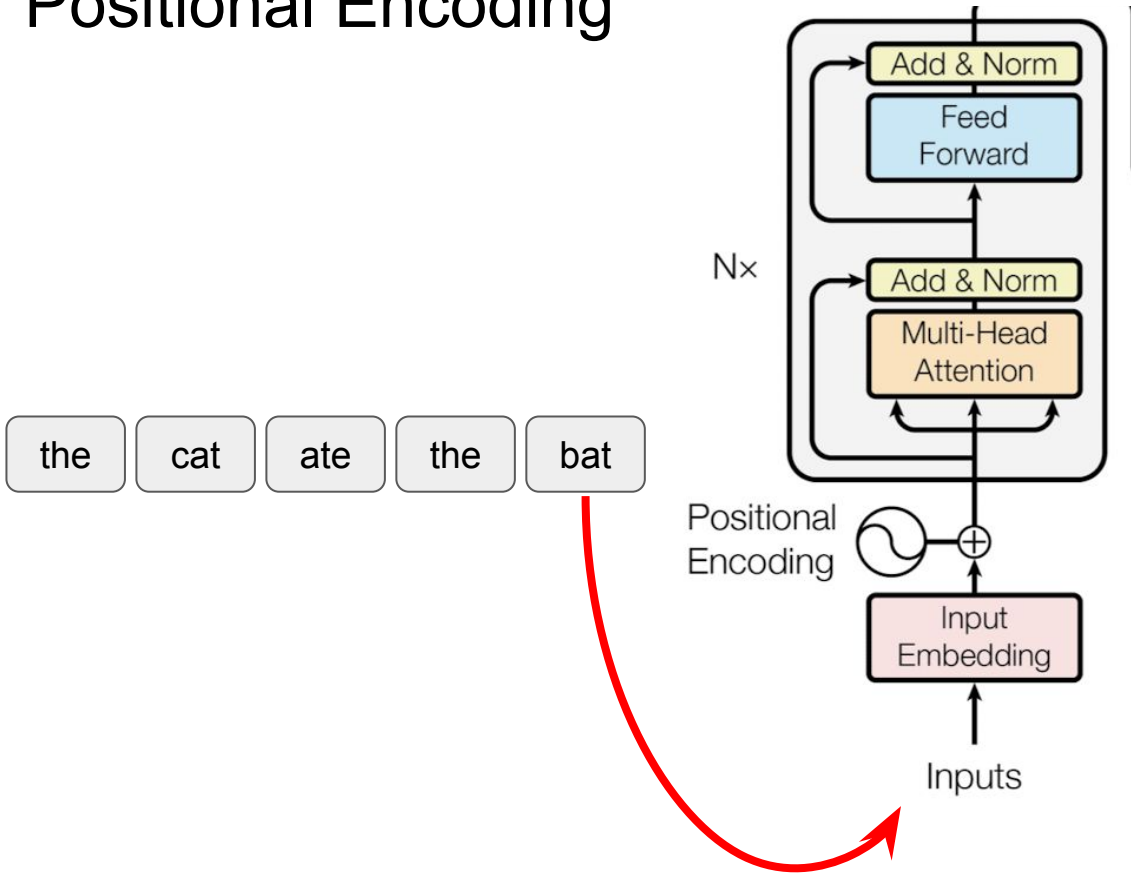

Positional Encoding

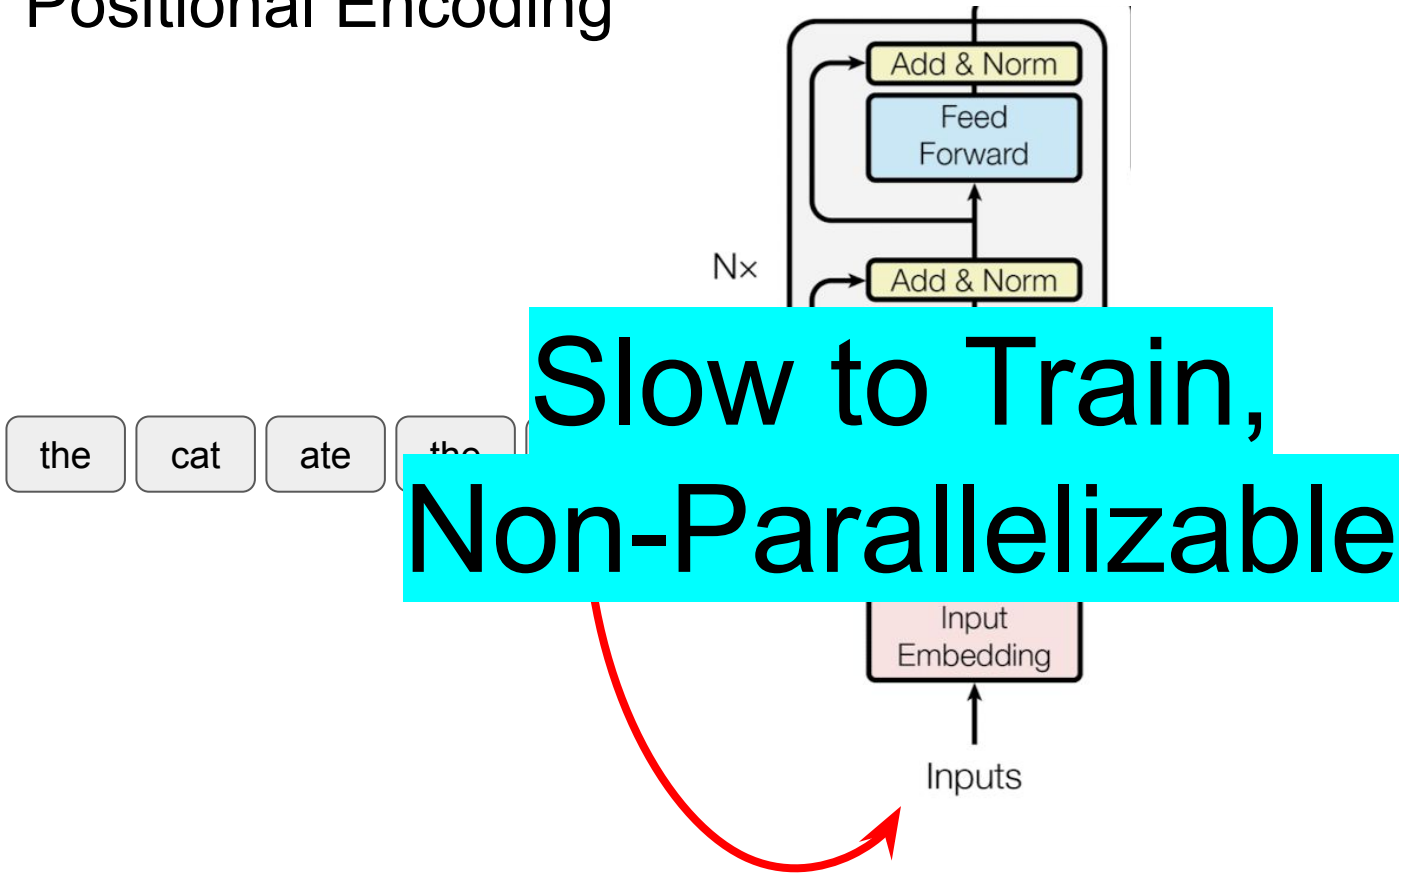

Slow to Train, Non-Parallelizable

Positional Encoding

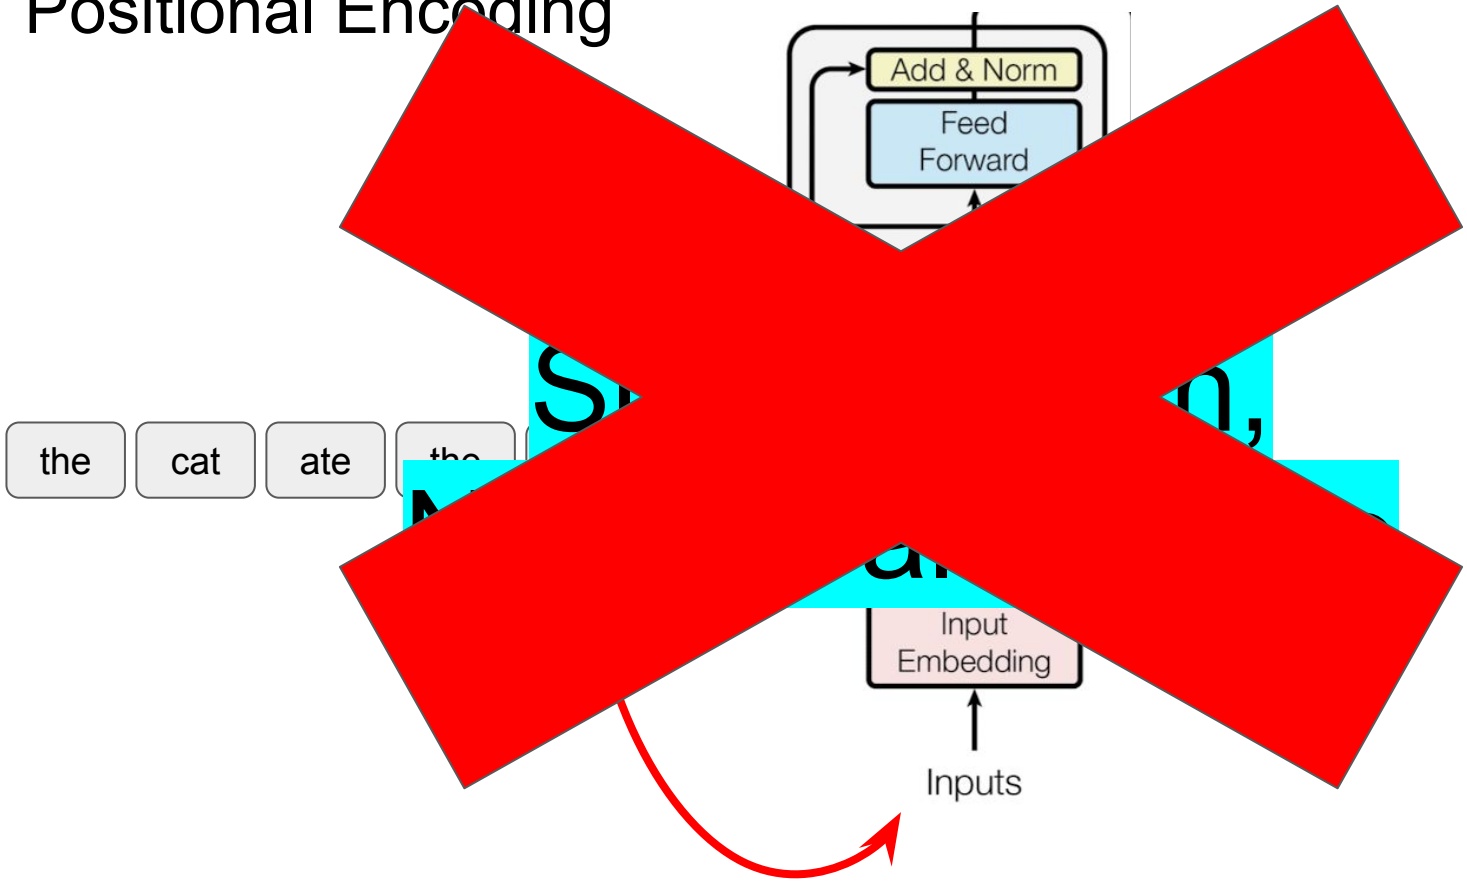

Positional Encoding

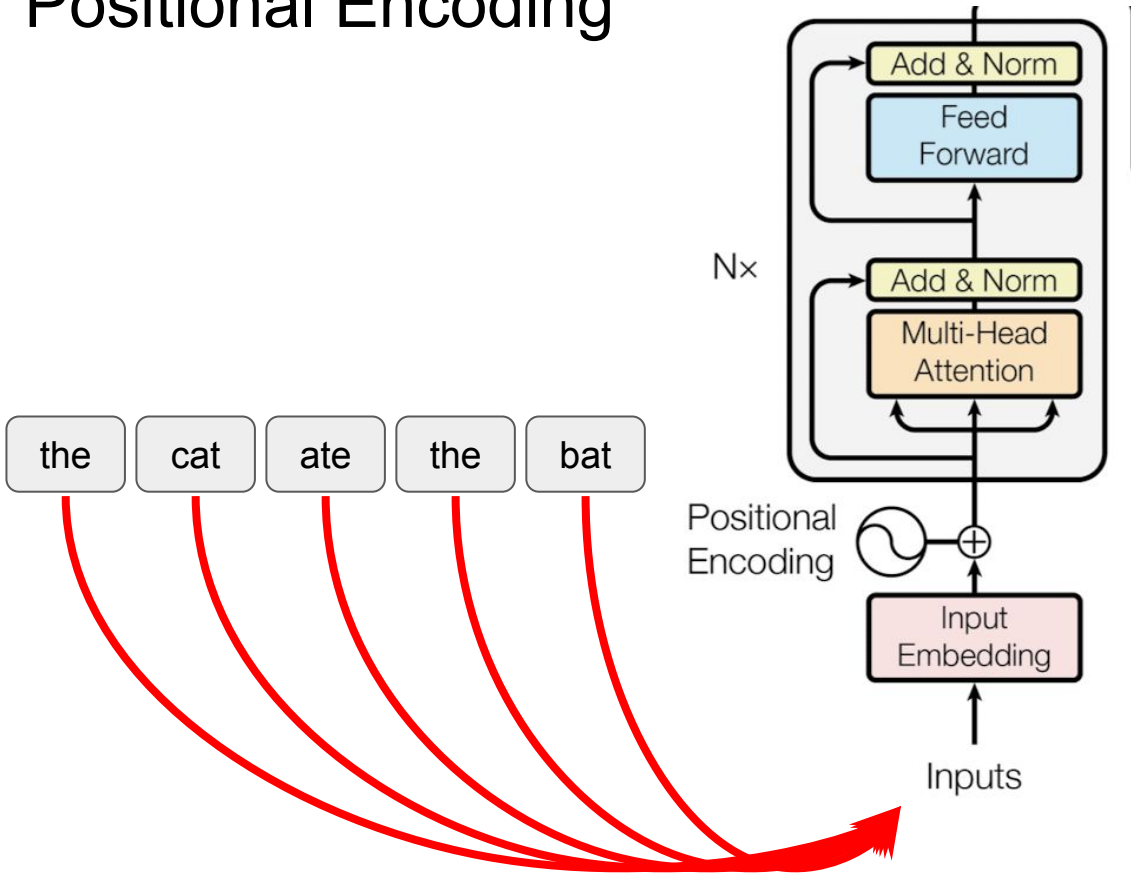

Word Embeddings + Positional Encoding

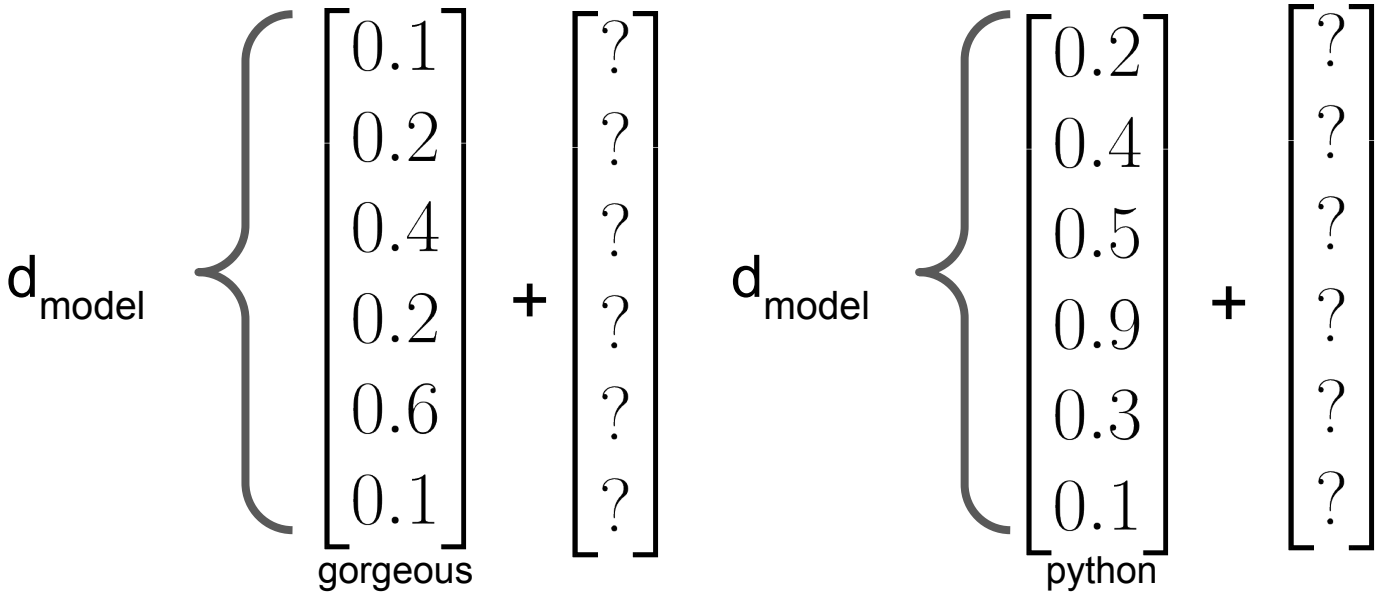

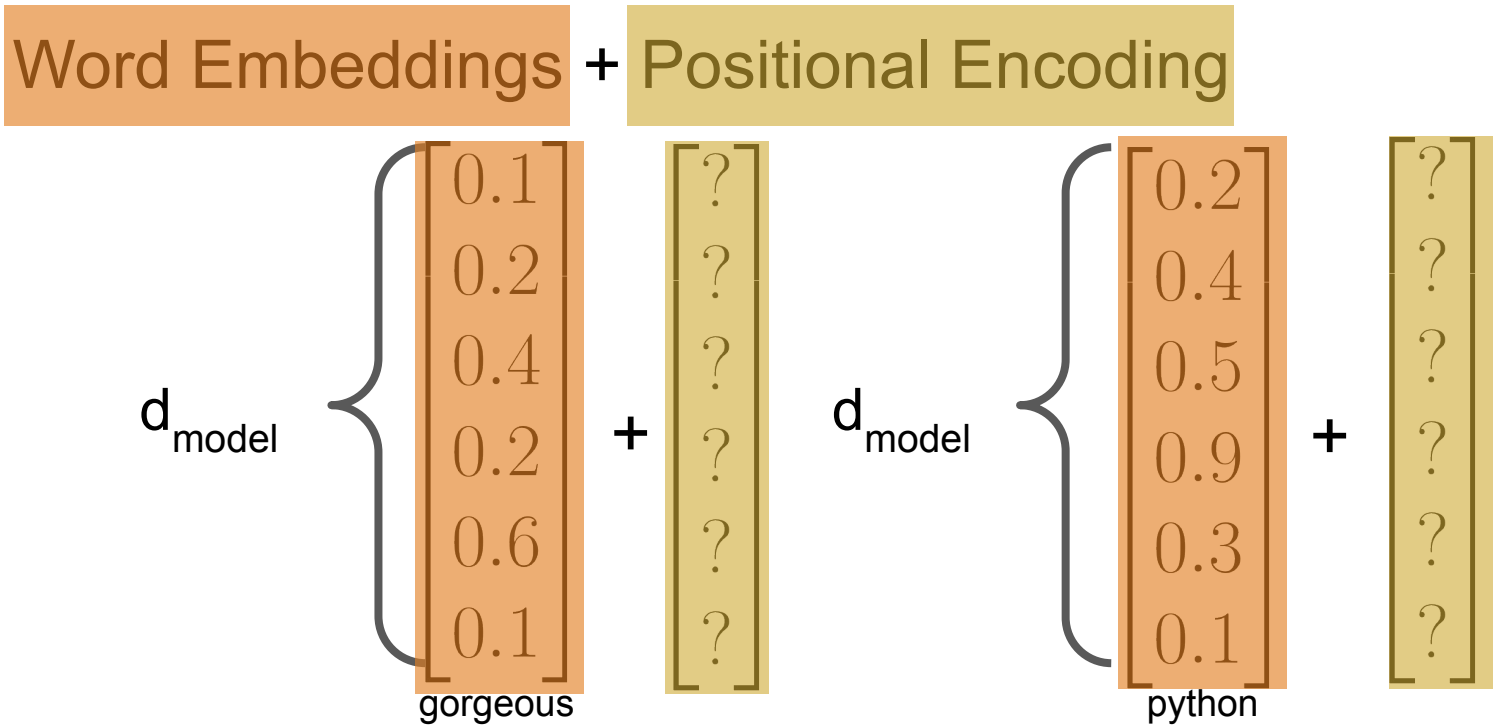
Word Embeddings
+
Positional Encoding
d model
0.1
0.2
0.4
0.2
0.6
0.1
?
?
?
?
?
?
+
gorgeous
d model
0.2
0.4
0.5
0.9
0.3
0.1
?
?
?
?
?
?
+
python

Generating Positional Encodings

Positional Encodings

your python code is gorgeous

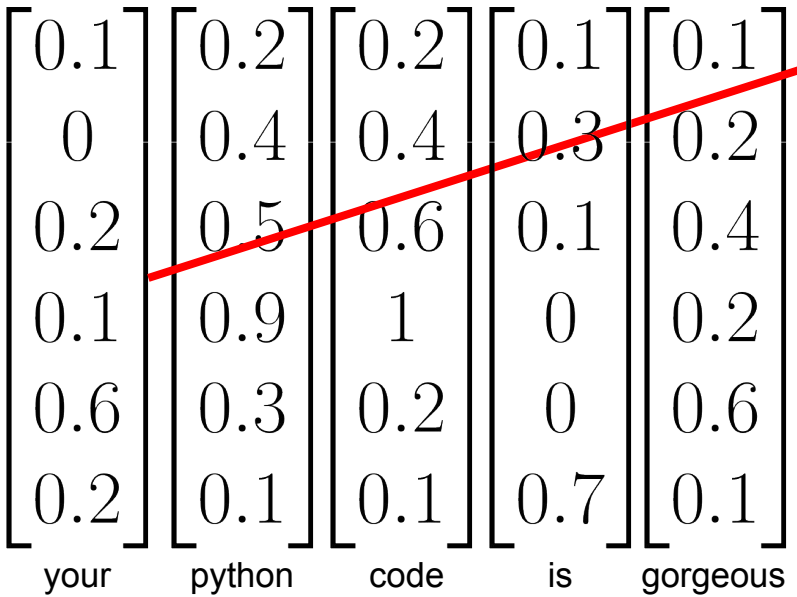

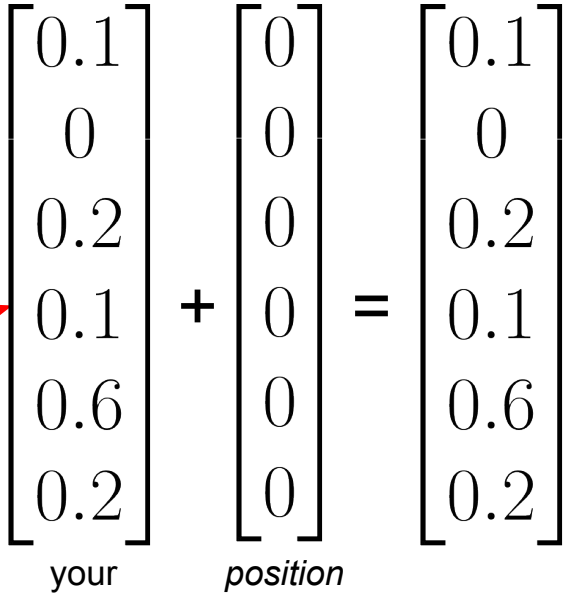

Positional Encodings

your python code is gorgeous

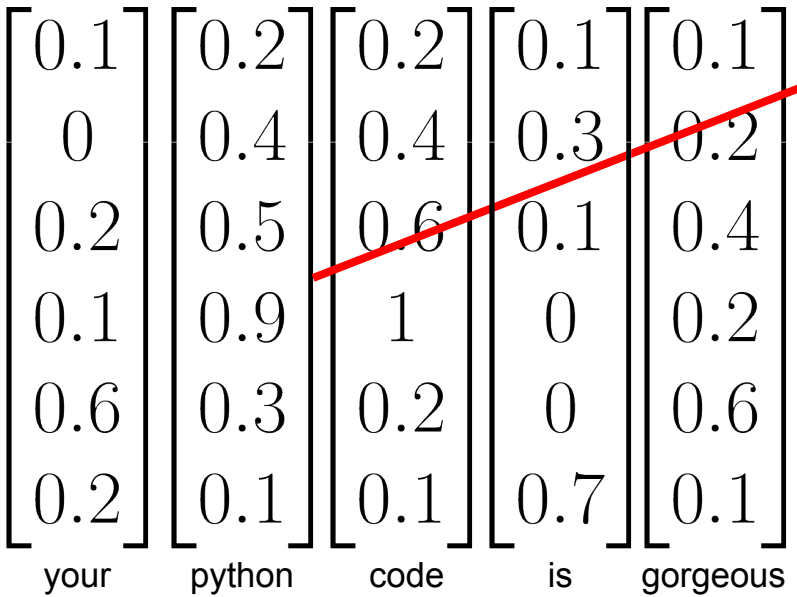

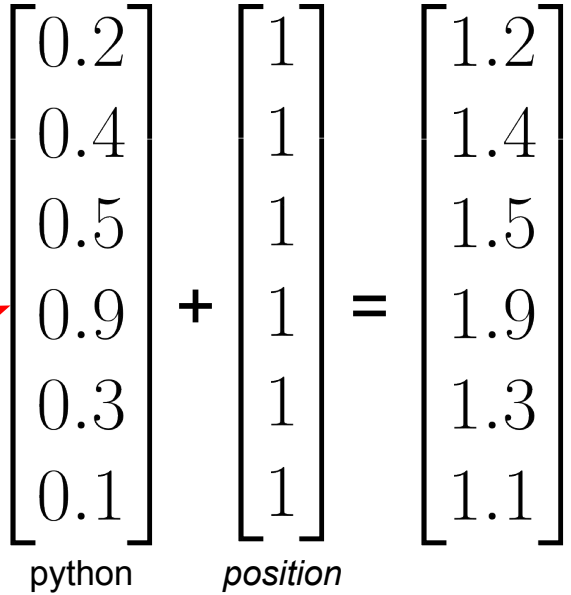

Positional Encodings

your python code is gorgeous

$$\begin{bmatrix}0.1\\0\\0.2\\0.1\\0.6\\0.2\end{bmatrix}\begin{bmatrix}0.2\\0.4\\0.5\\0.9\\0.3\\0.1\end{bmatrix}\begin{bmatrix}0.2\\0.4\\0.6\\1\\0.2\\0.1\end{bmatrix}\begin{bmatrix}0.1\\0.3\\0.1\\0\\0\\0.7\end{bmatrix}\begin{bmatrix}0.1\\0.2\\0.4\\0.2\\0.6\\0.1\end{bmatrix}$$

your python code is gorgeous

$$\begin{bmatrix}0.2\\0.4\\0.6\\1\\0.2\\0.1\end{bmatrix} + \begin{bmatrix}2\\2\\2\\2\\2\\2\end{bmatrix} = \begin{bmatrix}2.2\\2.4\\2.6\\3\\2.2\\2.1\end{bmatrix}$$

code *position*

Positional Encodings

Instead of:	0	1	2	3	4
Use:	0	0.25	0.5	0.75	1

Positional Encodings

your python code is gorgeous

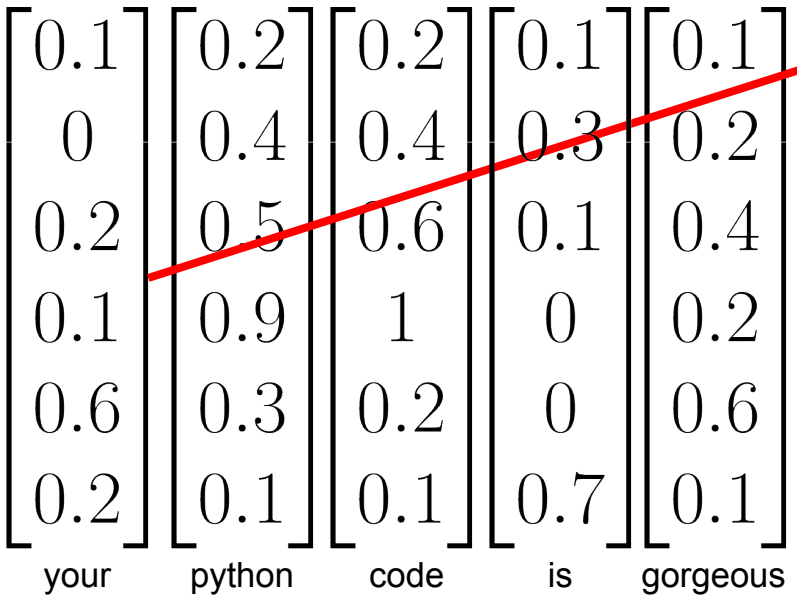

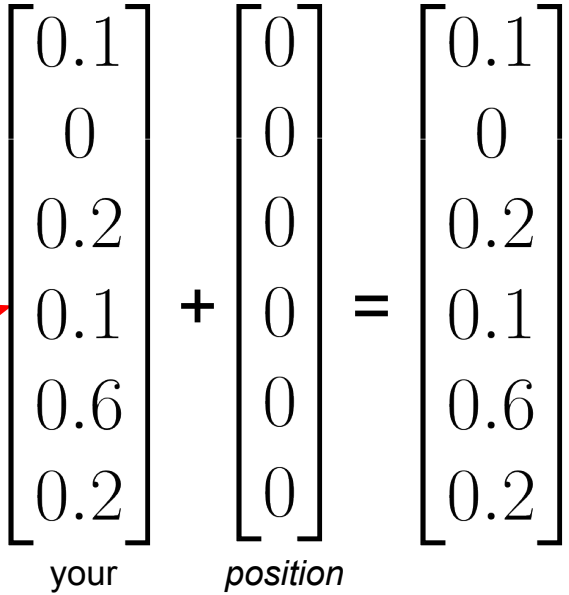

Positional Encodings

your python code is gorgeous

$$\begin{bmatrix}0.1\\0\\0.2\\0.1\\0.6\\0.2\end{bmatrix}\begin{bmatrix}0.2\\0.4\\0.5\\0.9\\0.3\\0.1\end{bmatrix}\begin{bmatrix}0.2\\0.4\\0.6\\1\\0.2\\0.1\end{bmatrix}\begin{bmatrix}0.1\\0.3\\0.1\\0\\0\\0.7\end{bmatrix}\begin{bmatrix}0.1\\0.2\\0.4\\0.2\\0.6\\0.1\end{bmatrix}$$

your python code is gorgeous

$$\begin{bmatrix}0.2\\0.4\\0.5\\0.9\\0.3\\0.1\end{bmatrix} + \begin{bmatrix}0.25\\0.25\\0.25\\0.25\\0.25\\0.25\end{bmatrix} = \begin{bmatrix}0.45\\0.65\\0.75\\1.15\\0.55\\0.35\end{bmatrix}$$

python *position*

Positional Encoding

Things we want:

- Encode position
- No large numbers
- Allow for variable sequence length

Positional Encodings

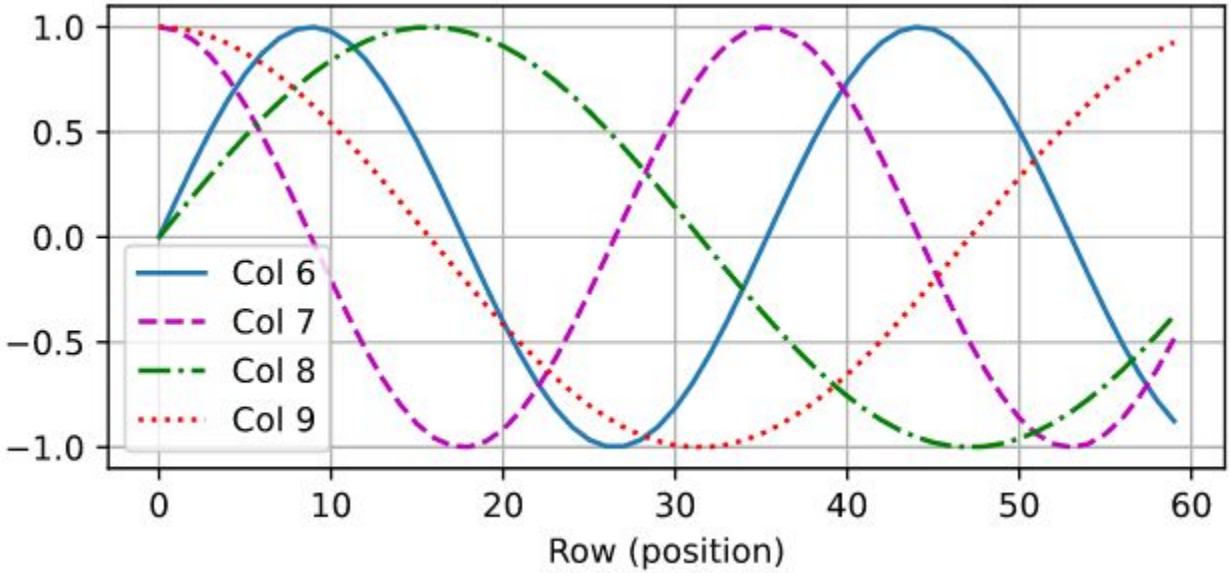

Positional Encoding

Pros

- sin() and cos() are bounded between -1, 1
- sin() and cos() can go on forever

Cons

- sin() and cos() are periodic

Positional Encoding

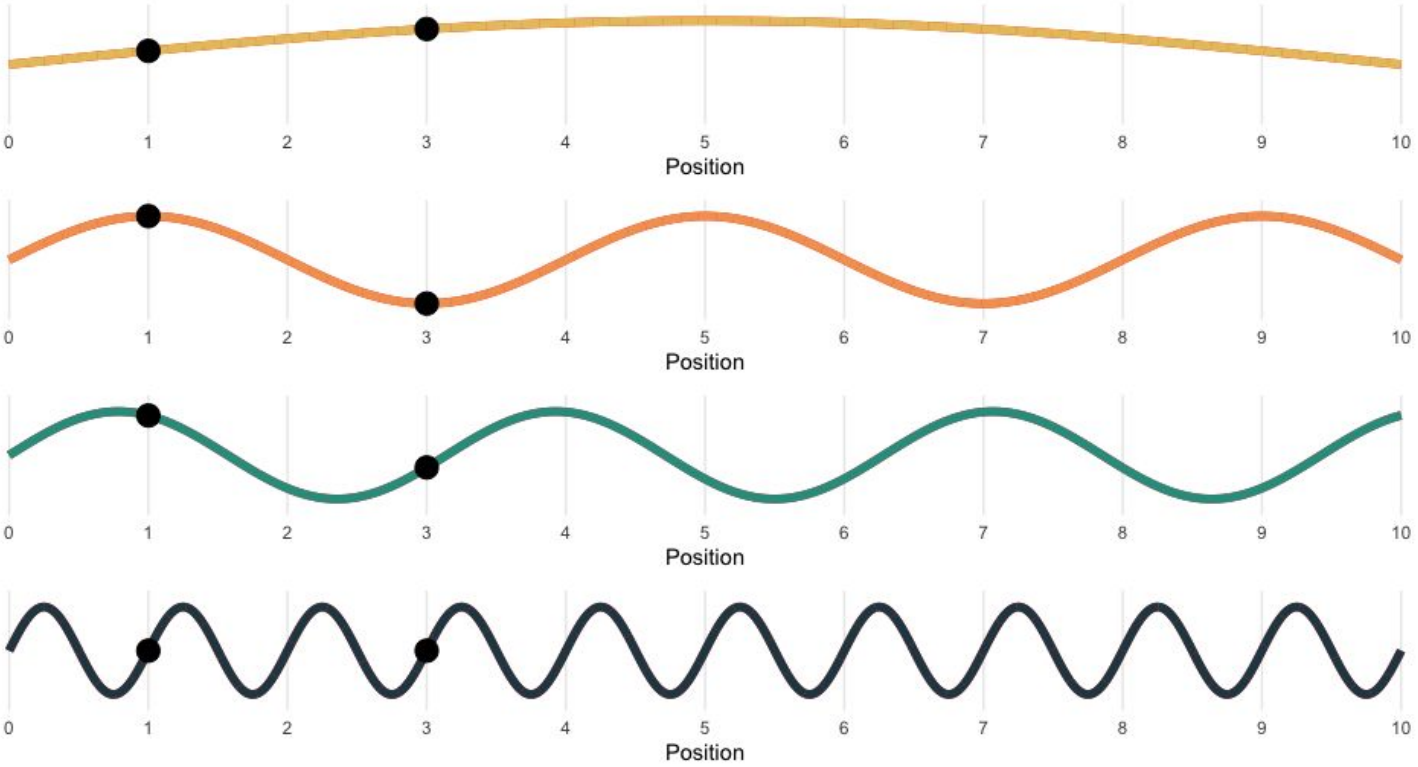

Positional Encoding

$$sin(pos/10000^{\frac{2*i}{d}})$$

$$cos(pos/10000^{\frac{2*i}{d}})$$

Positional Encoding

$$sin(pos/10000^{\frac{2*i}{d}})$$

$$cos(pos/10000^{\frac{2*i}{d}})$$

Positional Encoding

$$sin(pos/10000^{2*i/d})$$

$$cos(pos/10000^{2*i/d})$$

Positional Encoding

d_{model}

$$\begin{bmatrix} sin(pos/10000^{\frac{2*0}{d}}) \\ cos(pos/10000^{\frac{2*0}{d}}) \\ \\ sin(pos/10000^{\frac{2*1}{d}}) \\ cos(pos/10000^{\frac{2*1}{d}}) \\ \\ \dots \\ \\ sin(pos/10000^{\frac{2*j}{d}}) \\ cos(pos/10000^{\frac{2*j}{d}}) \end{bmatrix}$$

$$j = d_{model}//2$$

Positional Encoding

```
pos = 0

d = 6

pos_encoding = np.zeros(d)

for i in range(d//2):

    pos_encoding[2*i]     = sin(pos/10000^(2*i/d))

    pos_encoding[2*i + 1] = cos(pos/10000^(2*i/d))
```

$$\begin{bmatrix} sin(pos/10000^{\frac{2*0}{d}}) \\ cos(pos/10000^{\frac{2*0}{d}}) \\ \\ sin(pos/10000^{\frac{2*1}{d}}) \\ cos(pos/10000^{\frac{2*1}{d}}) \\ \\ \dots \\ \\ sin(pos/10000^{\frac{2*j}{d}}) \\ cos(pos/10000^{\frac{2*j}{d}}) \end{bmatrix}$$

$j = d_{model}//2$

Positional Encoding

```
pos = 0

d = 6

pos_encoding = np.zeros(d)

for i in range(d//2):

    pos_encoding[2*i]     = sin(pos/10000^(2*i/d))

    pos_encoding[2*i + 1] = cos(pos/10000^(2*i/d))
```

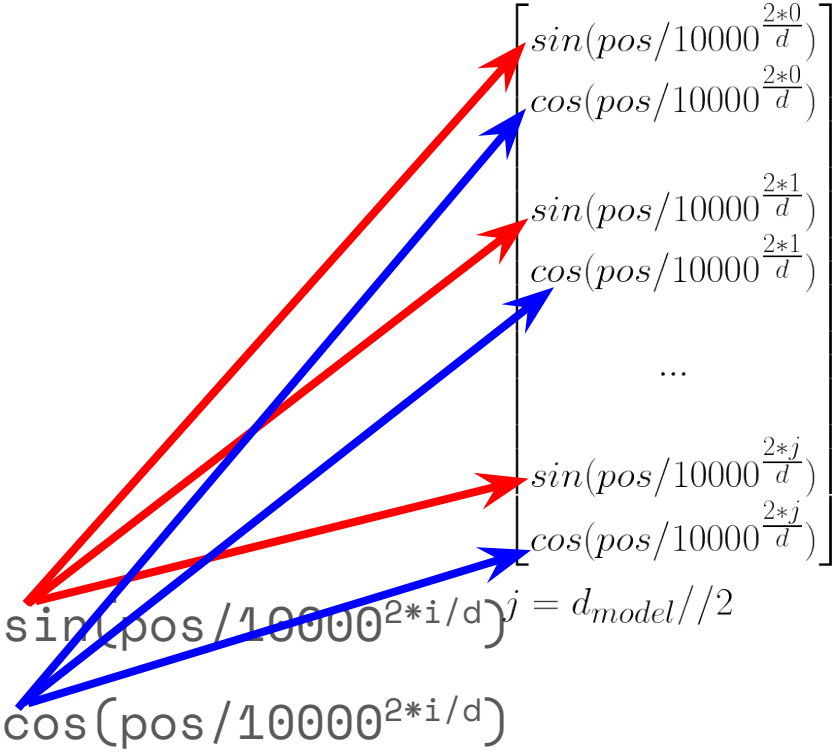

Positional Encoding

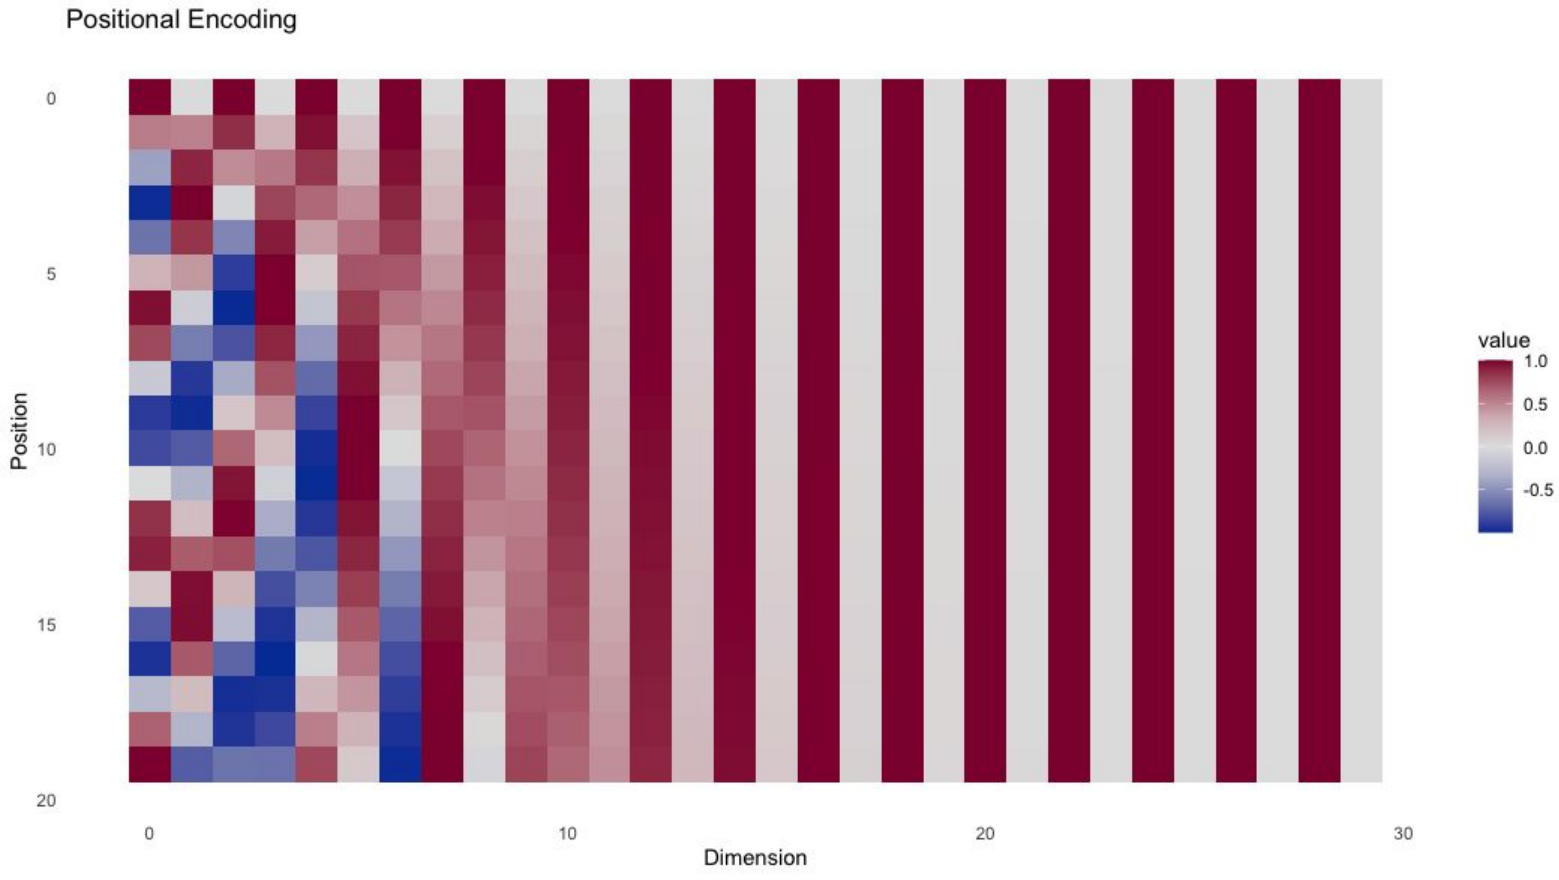

Positional Encoding

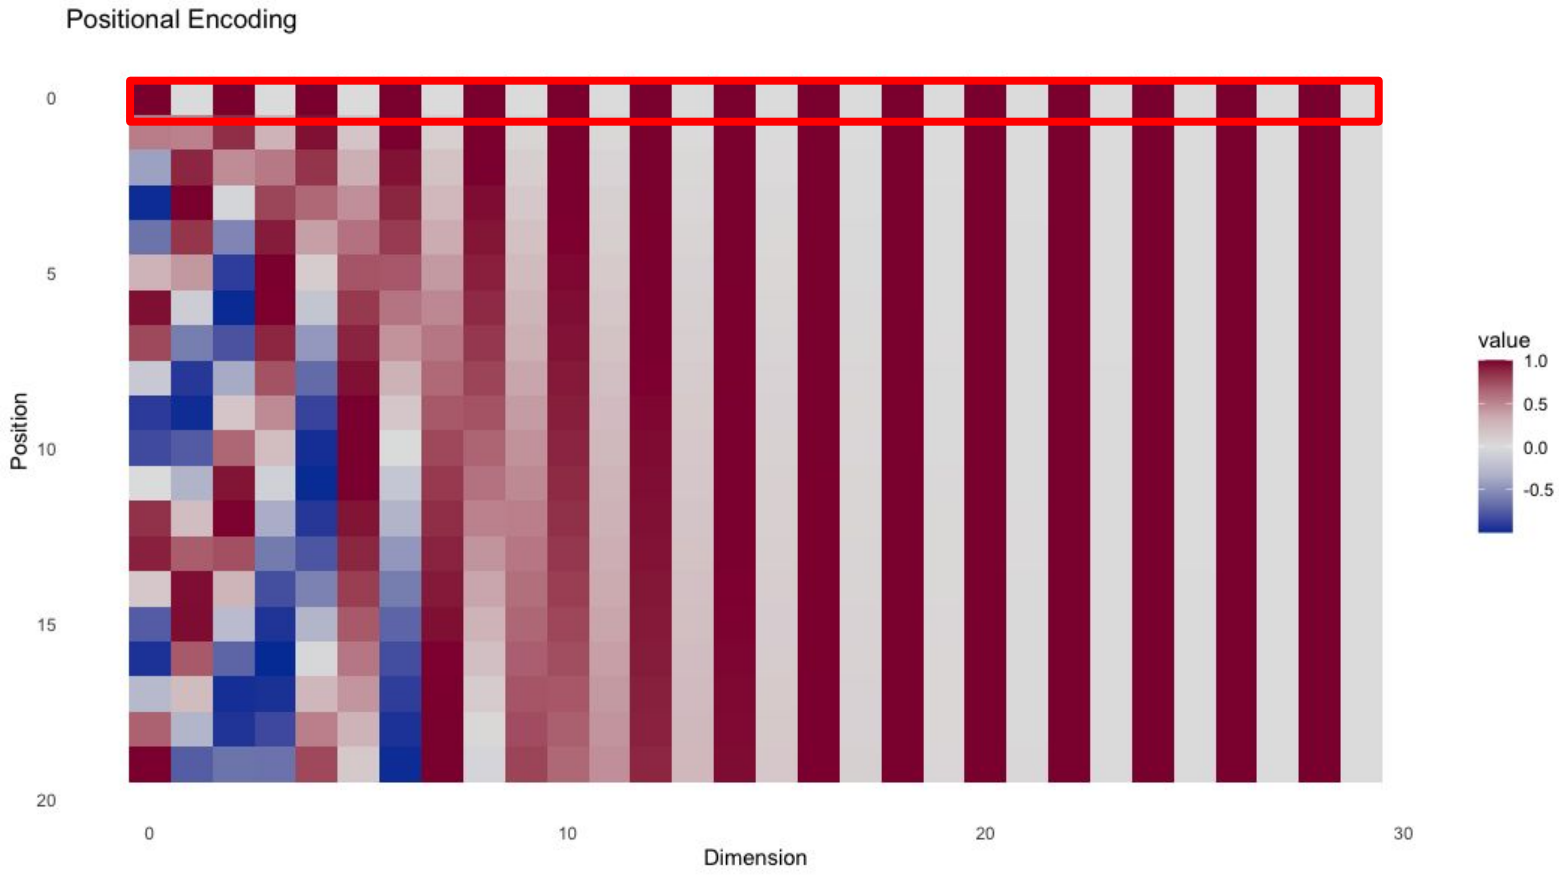

Positional Encoding

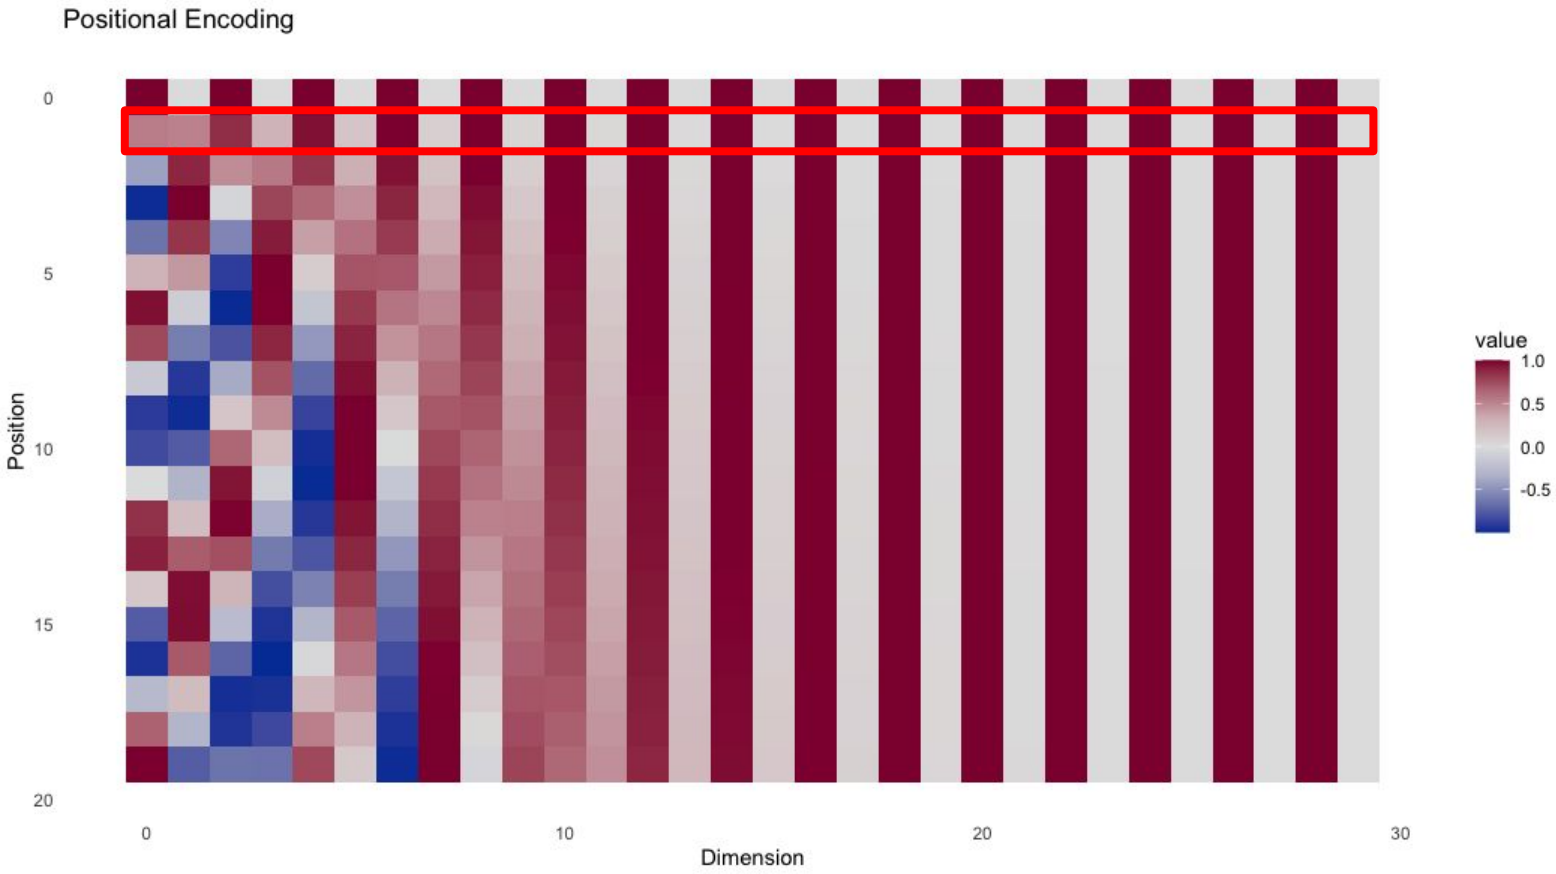

Positional Encoding

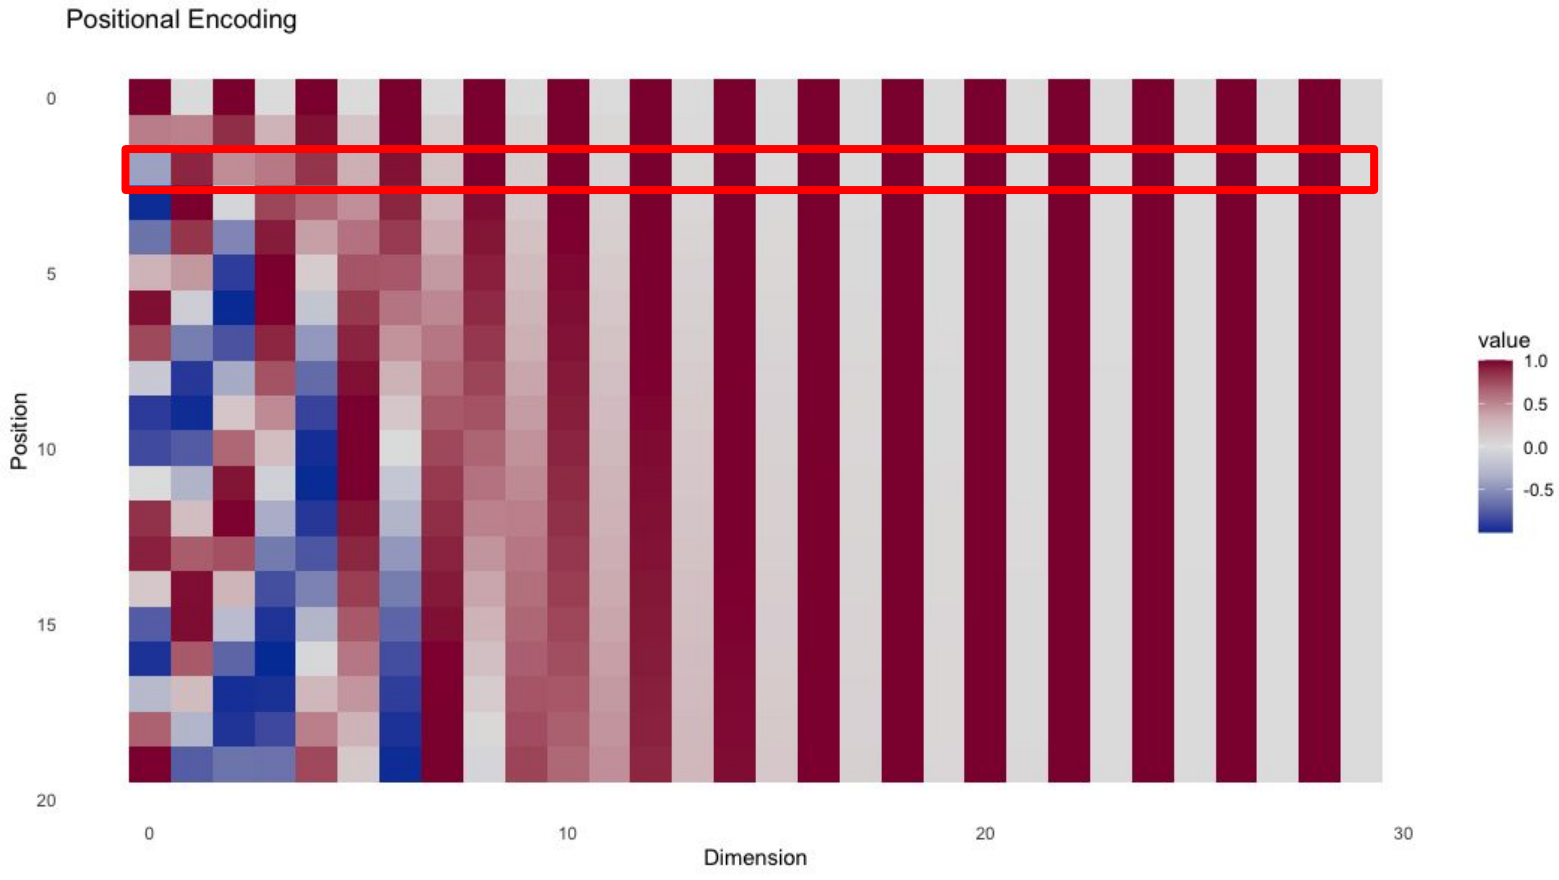

Positional Encoding

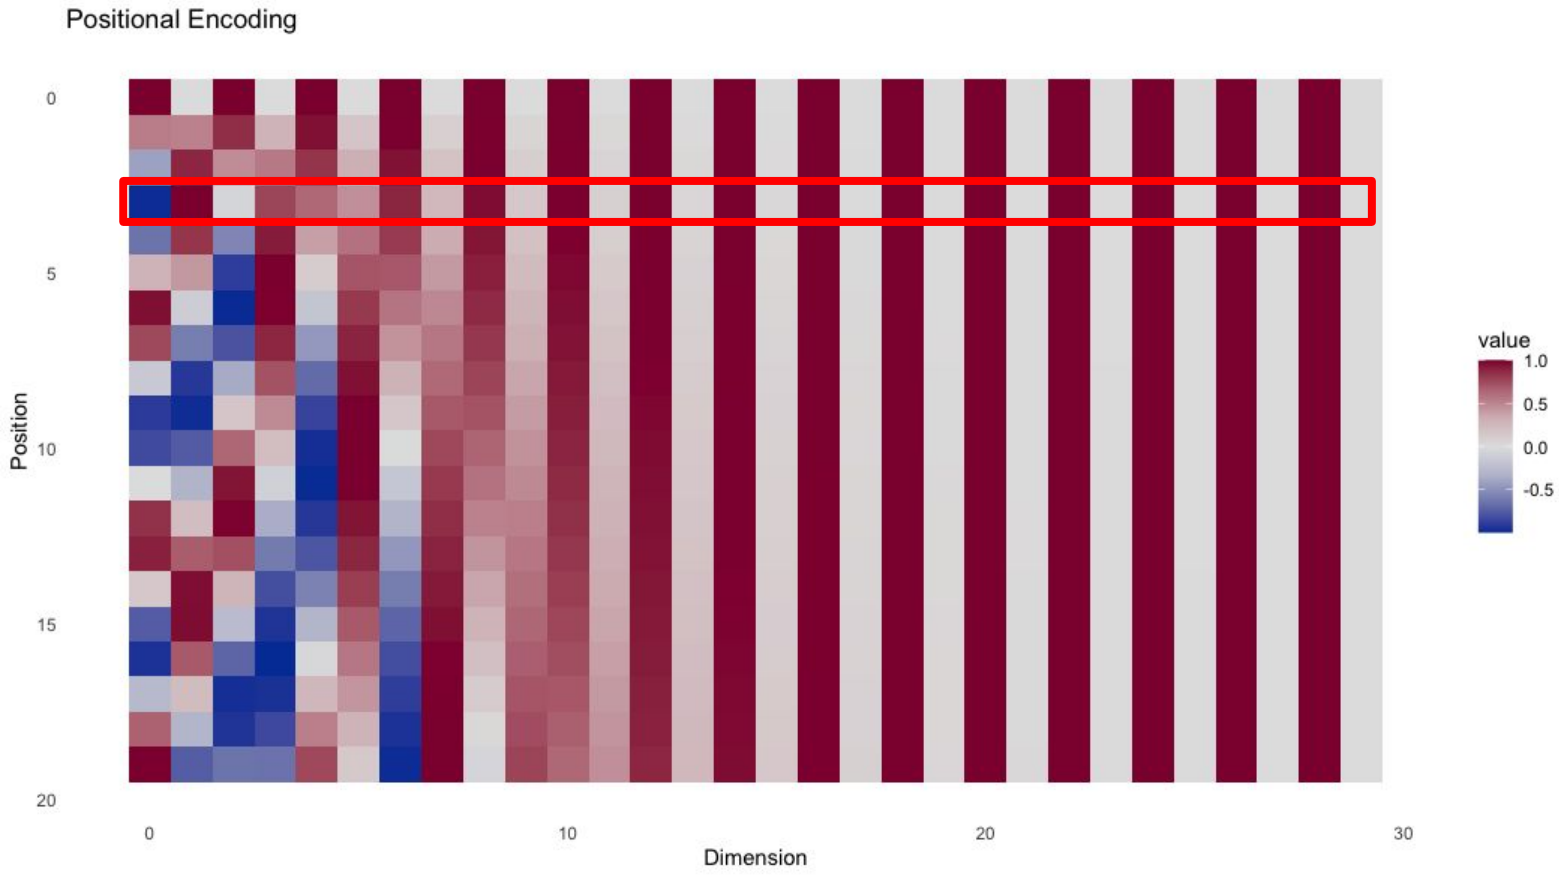

Positional Encoding

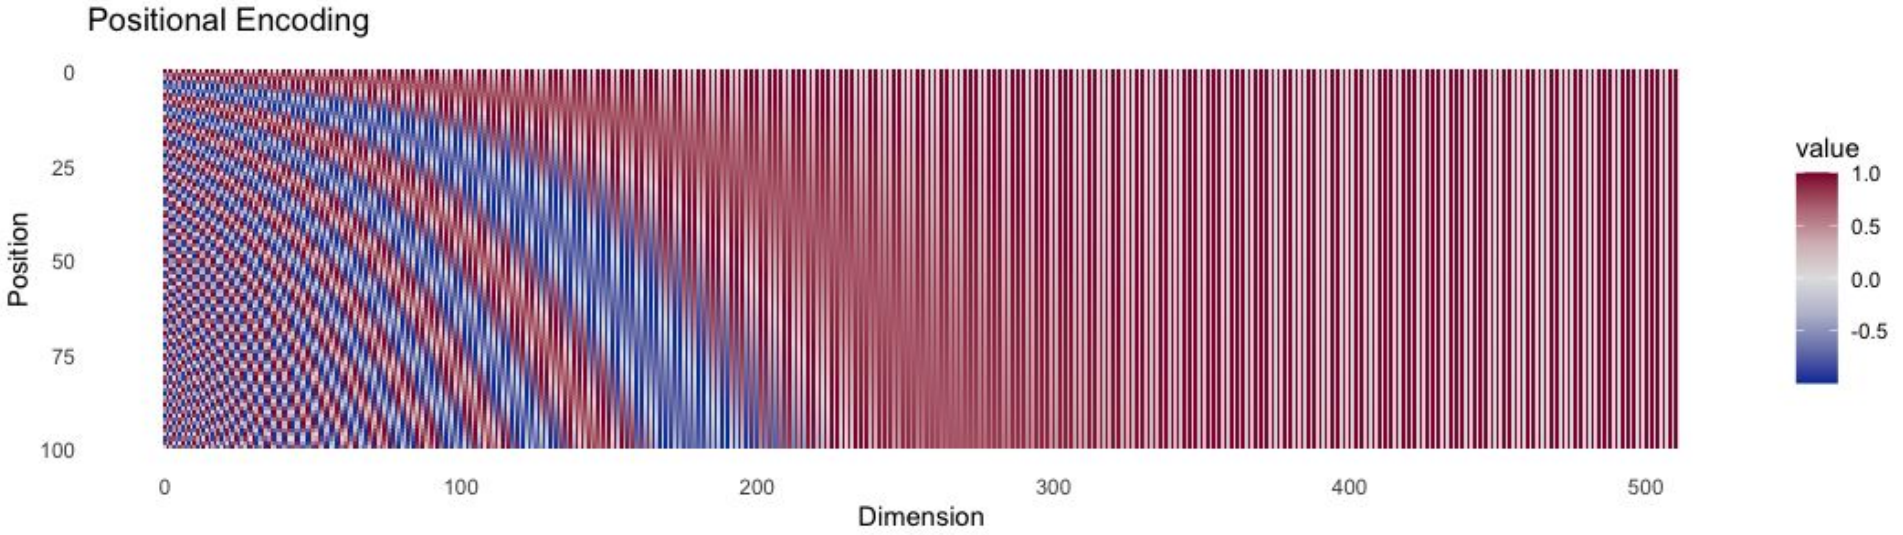

Positional Encodings

your python code is gorgeous

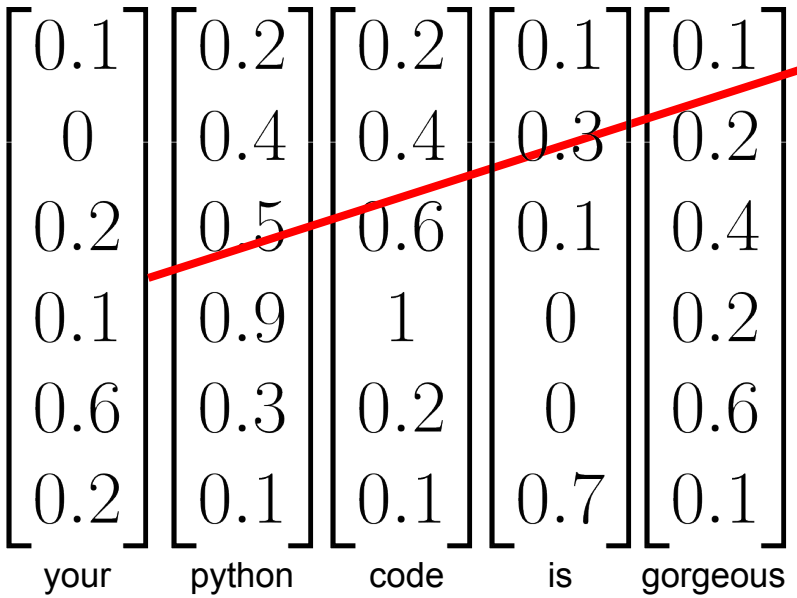

$$\begin{bmatrix} 0.1 \\ 0 \\ 0.2 \\ 0.1 \\ 0.6 \\ 0.2 \end{bmatrix} + \begin{bmatrix} sin(0/10000^{\frac{2*1}{6}}) \\ cos(0/10000^{\frac{2*1}{6}}) \\ sin(0/10000^{\frac{2*2}{6}}) \\ cos(0/10000^{\frac{2*2}{6}}) \\ sin(0/10000^{\frac{2*3}{6}}) \\ cos(0/10000^{\frac{2*3}{6}}) \end{bmatrix}$$

your *position*

Positional Encodings

your python code is gorgeous

$$\begin{bmatrix}0.1\\0\\0.2\\0.1\\0.6\\0.2\end{bmatrix}\begin{bmatrix}0.2\\0.4\\0.5\\0.9\\0.3\\0.1\end{bmatrix}\begin{bmatrix}0.2\\0.4\\0.6\\1\\0.2\\0.1\end{bmatrix}\begin{bmatrix}0.1\\0.3\\0.1\\0\\0\\0.7\end{bmatrix}\begin{bmatrix}0.1\\0.2\\0.4\\0.2\\0.6\\0.1\end{bmatrix}$$

your python code is gorgeous

$$\begin{bmatrix}0.2\\0.4\\0.5\\0.9\\0.3\\0.1\end{bmatrix} + \begin{bmatrix}sin(1/10000^{\frac{2*1}{6}})\\cos(1/10000^{\frac{2*1}{6}})\\sin(1/10000^{\frac{2*2}{6}})\\cos(1/10000^{\frac{2*2}{6}})\\sin(1/10000^{\frac{2*3}{6}})\\cos(1/10000^{\frac{2*3}{6}})\end{bmatrix}$$

python *position*

Parallelization

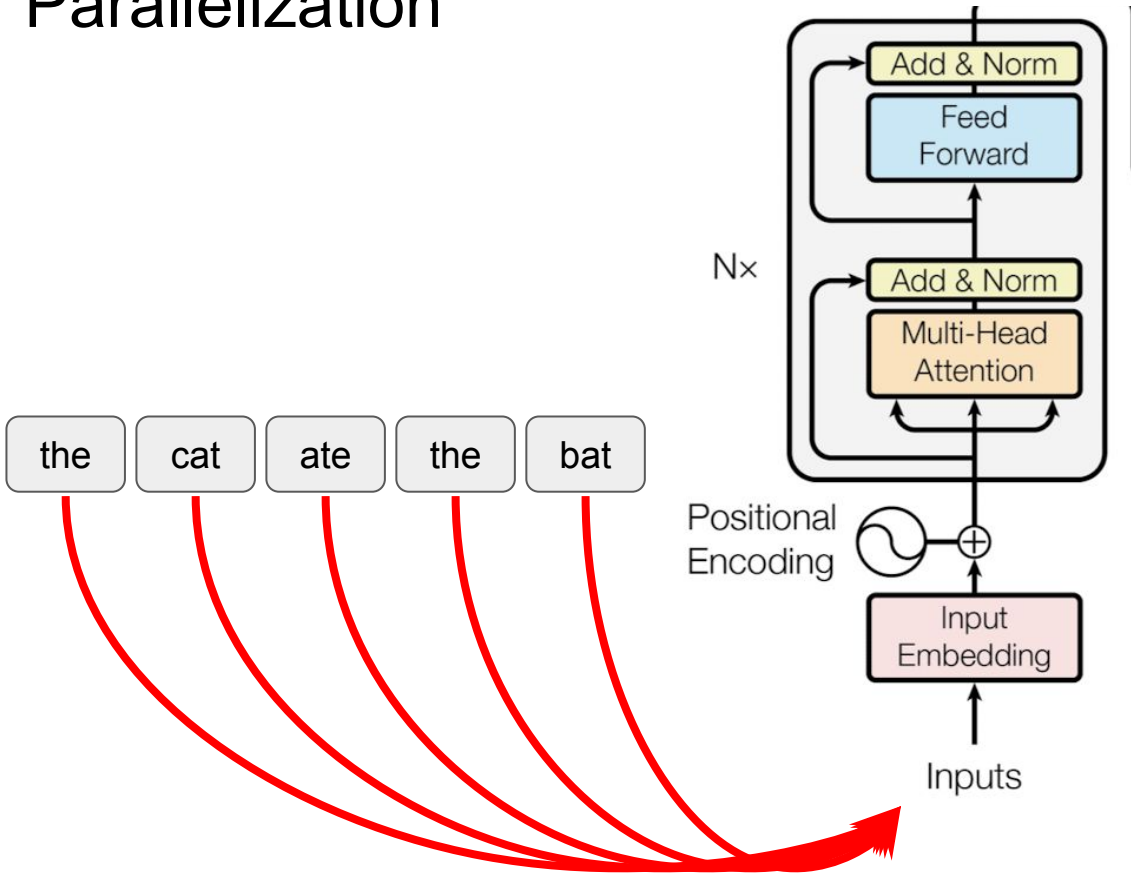

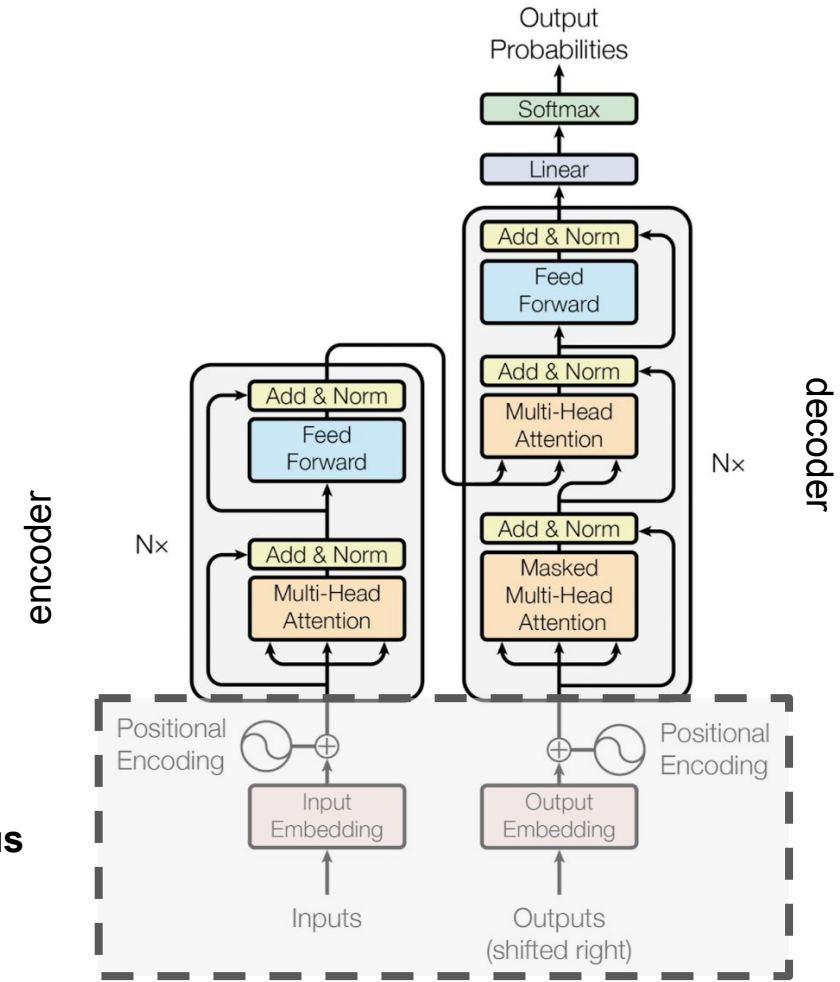

Figure 1: The Transformer - model architecture.

Today's Focus

Image from: Attention is all you Need (Vaswani et. al 2017)

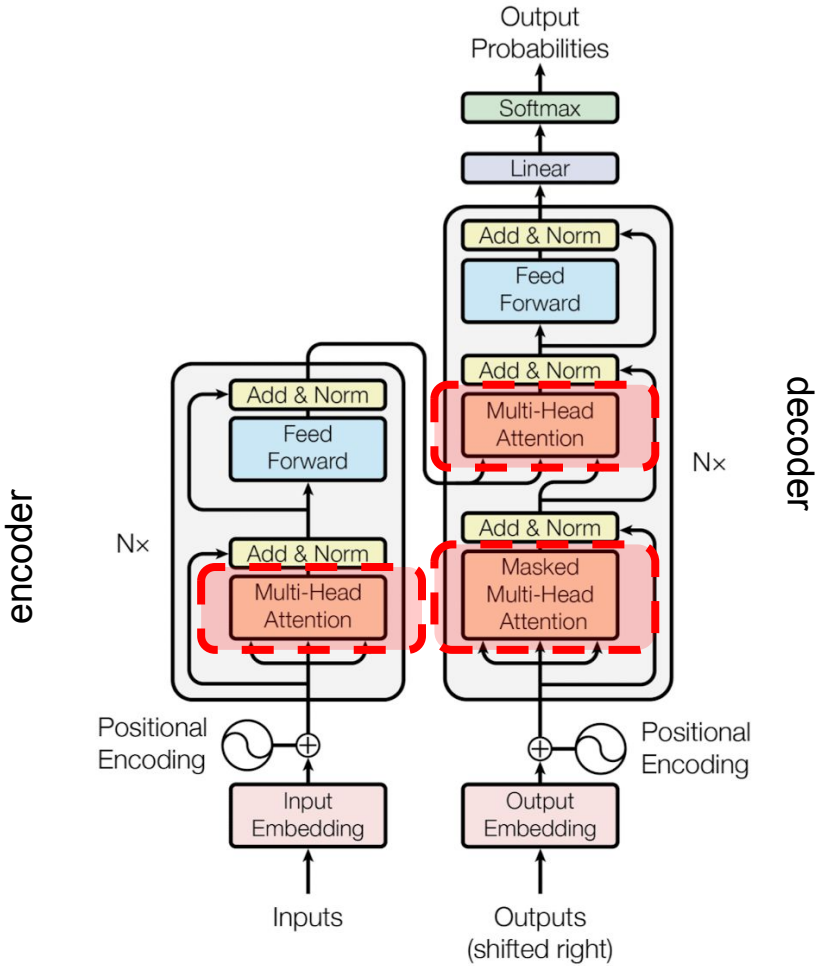

Figure 1: The Transformer - model architecture.

Image from: Attention is all you Need (Vaswani et. al 2017)